

DECISION FRAMEWORK · Q3 2025

From Signal to Strategy

A 78-slide answer to the question "have you considered writing things down"

What this deck covers, in order

01 Why This Exists — slide 3

→ **The Framework — slide 7**

03 The Evaluation — slide 22

04 The Build — slide 31

05 The Results — slide 47

CONTEXT • COMPETITIVE DYNAMICS

The window for differentiation is narrowing

Commoditized infrastructure, compressed release cycles, and rising switching costs have cut the average durable advantage from 36 months to under 14. This slide will appear in every deck for the next two years regardless of whether that stays true.

“

The signal was always there. We just didn't have a system that forced us to look at it before we'd already decided.

”

— Head of Product, Pilot Team 3, in a retrospective where we then decided what we'd already decided

IMPACT • PILOT RESULTS

Six months of results across four product teams

SAME TEAMS, SAME CONDITIONS, SAME SPREADSHEET.



CALIBRATION RESULT • 6-MONTH PILOT

14X

Return on signal investment — on a baseline
the framework team defined. Verified by
nobody.

SECTION 01 • THE FRAMEWORK

We built a four-component
scoring system. Two of the
four are used regularly


SIGNAL DEFINITION • WORKSHOP 04

Before we score signals, we need to agree on what a signal is

Three workshops in. The definition gets
socialized in a fifth before anyone can use it.

How signals move from input to decision

Four-stage pipeline — 11 weeks in, still in pilot



The framework has four components

Signal Intake

Weekly structured collection across customer conversations, market data, and competitive moves. Everyone agrees it's a good idea. Nobody does it the week of the retrospective.

Scoring Model

Confidence, recency, and strategic relevance, each weighted. The weights are team-configurable — which means the head of product reconfigures them until the output agrees with the roadmap.

Decision Log

Every decision recorded with its signals and criteria. A required artifact. It holds 18 entries, against roughly 340 decisions actually made.

Calibration Loop

A monthly retrospective comparing predicted outcomes to real ones. The meeting reliably exists. The predictions, reliably, do not.

Signal Intake produces three outputs

1

Weekly Signal Brief

Last week's top 10 signals, ranked, with confidence scores and sources. Sent to product leads every Monday — where it sits unread in a folder called "Framework Stuff."

2

Anomaly Alerts

Real-time flags when a signal crosses the 2σ threshold, routed to the accountable PM on a 4-hour SLA. The PM usually replies inside the 4 hours — to ask what 2σ means.

3

Monthly Signal Index

The source of truth for the calibration loop, and required reading before every retrospective. Nobody has read it. It is comprehensive.

The three things the framework connects

Signal

The observed input — a verbatim, a metric move, a competitor's announcement. The unit of intake. Frequently confused with "things the VP heard at a conference."

Decision

A signal plus a deadline, logged with its rationale. Every decision is meant to trace to a signal. In practice the signal is "we discussed it at the offsite."

Outcome

The result, measured against the prediction at retrospective. The unit of calibration. 18 have been logged. Roughly 340 have occurred.

Two failure modes the framework is designed to prevent

False signal amplification.

A single loud voice — one enterprise customer, one analyst report, one VP with a feeling — dominates the decision without ever being weighed against the full signal set. The model caps any one source at 30% of total weight. Unless that source is the CEO, in which case the cap is a guideline.

Signal hoarding.

Teams collect signals but never log decisions, so the calibration loop has nothing to learn from. The rule — no log, no change above P2 — was printed on a poster and hung in the meeting room. The poster has since been replaced by a free-pizza flyer.



Scoring Model Deep Dive

What the scoring model actually does

The most configurable component — a feature or a warning sign, depending on your team. Three dimensions:

- 1

Confidence
Independent corroborating sources, 1–5. Enterprise customers count as 1, regardless of volume.
- 2

Recency
Time-decay, configurable half-life. Set to two weeks, then surprised only recent news scores.
- 3

Strategic Relevance
Owner-scored, 1–5. A 5 correlates with whoever is presenting the roadmap this quarter.

The six signal dimensions, what they measure, and how they are scored

01	Confidence	Independent sources corroborating the signal	1-5 • AUTO • ENTERPRISE ALWAYS GETS A 4
02	Recency	Time-decay from signal date, configurable half-life	0.0-1.0 • AUTO • "SHORT" AFTER A BAD QUARTER
03	Relevance	Alignment to current strategic bets, owner-scored	1-5 • MANUAL • WHAT THE PM ALREADY PLANNED
04	Reach	Customers or segments affected	1-5 • AUTO • 5 IF THE ENTERPRISE ASKED
05	Effort	Engineering and design cost to act on the signal	1-5 • MANUAL • ENGINEERING'S EYEBROWS GO UP
06	Confidence delta	Change in confidence since the last scoring cycle	-5 TO +5 • AUTO • DID ANYONE TALK TO A CUSTOMER

Scoring model: before and after the calibration loop

BEFORE CALIBRATION

Equal weights — confidence, recency, and relevance each contribute 33%. Simple, and at least honest that we are basically guessing.



AFTER CALIBRATION

Weights reflect your team's historical accuracy — except the team keeps changing, and the history is three months of data from a quarter everyone agrees was atypical.

The shift from equal to calibrated weights takes two retrospective cycles — 60 days from adoption, or 14 months, depending on who you ask.

Two intake modes for different signal types

Structured Intake

Clear-schema signals — NPS verbatims, support tickets, win/loss notes. Ingested and scored automatically, with zero manual handling. Produces 94% of the data and 12% of the roadmap decisions.

Unstructured Intake

No-schema signals — field notes, conference hallway talk, a board member who had thoughts at a dinner. Routed to the owner for manual classification. Produces 6% of the data and 88% of the roadmap decisions.

How a decision moves through the framework



What the framework does not do

Doesn't make decisions — it formalizes the one you'd have made anyway.

Doesn't replace discovery — it routes whatever surfaces before the roadmap locks.

Doesn't work without the Decision Log. 18 entries; ~340 decisions made.

Doesn't guarantee alignment — the same argument, just earlier in the quarter.

Doesn't scale below P2 — the daily 90% still lives in Slack threads from 2022.

Four things that must be true before you begin

- 01 A monthly prioritization cadence. "Whenever someone escalates loudly" doesn't count — it just gets louder.
- 02 Someone who owns signal collection. Which means defining a signal first. See slide 8.
- 03 Leadership logging rationale, not just outcomes. This slide exists because that part is hard.
- 04 90 minutes a week. Budget four.

SECTION 02 WEIGHS FOUR TOOLS AGAINST FOUR
CRITERIA — DEFINED BY THE TEAM THAT BUILT ONE.
DISCLOSED IN A FOOTNOTE.

SECTION 01 OF 05 COMPLETE

The framework is
specified — the real
question is build or buy

SECTION 02 • THE EVALUATION

We evaluated four tools.
One of them was built by
the evaluation team

Four requirements every decision system must meet

01

Speed

Decisions close inside their window. Six months to calibrate isn't fast. We won't dwell on that here.

02

Auditability

Every decision above the threshold carries a traceable rationale — for after the launch goes badly.

03

Adoption

Weekly use, or calibration never runs. We budgeted 90 minutes per PM. Actual: 11.

04

Calibration

It has to improve. A static model is a spreadsheet with a dashboard nobody checks.

We evaluated four intake tools against the criteria

Tool A • Chorus

- ✔ Speed
- ✖ Auditability
- ✔ Adoption
- ✖ Calibration

Great call recording, no logging, no calibration. The sales team already uses it.

Tool B • Productboard

- ✖ Speed
- ✔ Auditability
- ✔ Adoption
- ✖ Calibration

Solid intake, no calibration. Setup was "3–4 weeks." It took 11.

Tool C • Notion

- ✔ Speed
- ✔ Auditability
- ✖ Adoption
- ✖ Calibration

Build anything. Teams built seven, then debated the canon for two quarters.

Tool D • Sprig + Decision Log



- ✔ Speed
- ✔ Auditability
- ✔ Adoption
- ✔ Calibration

Built by this evaluation's authors. Meets all four — as they defined them.

The four tools side by side

CRITERION	CHORUS	PRODUCTBOARD	NOTION	SPRIG + LOG
Speed	✓	✗	✓	✓
Auditability	✗	✓	✓	✓
Adoption	✓	✓	✗	✓
Calibration	✗	✗	✗	✓
Setup time	1 day	3–4 weeks	40+ hours	Same day

♦ Criteria defined by the team building Sprig + Log. We're transparent about it. It's in the footnote.

How we sort the four tools against our two axes

COVERAGE • COST

High coverage / Low cost

Sprig + Log — best coverage, lowest TCO. Built by the evaluation team.
Productboard — narrower, here to show we looked.

High coverage / High cost

Notion — full coverage in theory. We tried it: seven versions, a quarterly canon debate.

Low coverage / Low cost

Chorus — cheap, three criteria short. The sales team likes it.

Low coverage / High cost

None — either the signal or a gap. We're treating it as a signal.

Applying the criteria to the tools — here is where the evidence points

The evidence favors Tool D



Sprig plus a lightweight Decision Log meets all four criteria inside the 90-minute weekly budget, reaches production the week it's adopted, and was recommended before the evaluation began. The evaluation confirmed this.


The path is not self-executing

Sprig needs a connector to your NPS and support platforms — budget 4–6 hours in week one. The connector has been in review for 11 weeks. We still call it "week one."

The Decision Log is the hardest part

Not technically — culturally. PMs have to log decisions before they close, not after. It's been described as a habit change in every deck since Q3 2023. It remains one.

The evaluation came down to one question the vendors could not answer



The left card is struck through; the DECISION connector is bold. The conclusion came first — the slide was built to hold it.

We are building, not buying

DECISION • 2026 Q1

BUILD

Owens the scoring policy, the calibration loop, and the timeline — plus the maintenance, the on-call rotation, and every future explanation of why the framework scored the wrong thing.

WHY NOT BUY

Three vendors, none exposing calibration weights. The weights are the product; you cannot buy the product without them. That was the finding.

WHY NOT DELAY

The window closes in 18 months — a sentence that has been in this deck, unchanged, since Q1 2025.

How we make calls when the spec is silent

- 01 Default to the choice that's cheaper to reverse — unless reversing it needs a meeting.

- 02 Name the actor, never the system. The system can't be held accountable. The actor can be reorganized.

- 03 Write the bet on the slide with the choice. Not always where it gets reviewed.

What the build covers

Three phases, four workstreams. We own the policy, the loop, the timeline — and whatever Phase 2 becomes.

1

Phase 01 — Architecture

Scoring policy live, the Decision Log accepting entries, one pilot team in weekly cadence.

2

Phase 02 — Operations

Multi-team calibration, automated weight updates, and before-after data you could defend in a board update.

3

Phase 03 — Scale

Org-wide enablement. In the roadmap since 2024. We remain committed.

T

SECTION 03 • THE BUILD

Three phases ship the architecture, the operations, and the scale — and Phase 01 has shipped

What ships in each phase, by workstream

WORKSTREAM	PHASE 01	PHASE 02	PHASE 03
Signal Intake	Connector v1	Multi-source dedupe	Anomaly auto-routing
Scoring	Equal-weights model	Per-team calibration	Per-decision profiles
Decision Log	Append-only schema	Outcome auto-pairing	Examiner export
Adoption	One pilot team	—	Org-wide enablement

Phase 3 holds org-wide enablement. In the roadmap since 2024. Phase 2 is ongoing.

Who owns each part of the framework lifecycle

Signal custody

Signal owner

Manages intake quality. Tunes weights only by choosing what to surface.

Policy

Framework operator

Owns scoring, cadence, rollback. One person. Noted in the risk register.

Consumption

Product team

Runs intake and logging. Mostly asks the operator to change the weights.

Oversight

Auditor

Reads the audit trail. Read it once, found 18 entries, asked if that was expected.

Three phases get us from decision to org-wide adoption

PHASE 01

Architecture

Scopes what we build, what we buy, and what we defer. The output is an architecture decision record — which will itself need an architecture phase before it can be approved.



PHASE 02

Pilot

One team, one decision type, one quarter. The phase ends at production cadence — meaning the retrospective happened at least once.



PHASE 03

Rollout

Five teams in two months, ending above 90% adoption. Planned for Q2 — as it has been for three consecutive years.

Three milestones mark Phase 01 complete

MILESTONE A

Scoring policy in production

The signed policy runs end-to-end and the first calibrated brief lands in leadership's inbox. Their first reply asks whether the weights can be adjusted. They can.

MILESTONE B

Per-team weights

One framework carries distinct weights per team without forks; recalibration is a single update. Each team will still want its own.

MILESTONE C

Per-decision-class profiles

Authoring a profile for one decision class takes minutes. Agreeing what counts as a decision class takes a workshop series. See slide 8.

Phase 01 acceptance — what shipped, what slipped, what stayed open

✓	Scoring policy live across all pilot teams	
✓	Decision Log audit trail readable by Auditor role	shipped 2026-Q1
✓	One reference team running weekly cadence	one team
—	Examiner pack auto-generation	was Phase 01 • now Phase 1b
○	Adoption above 90%	Phase 03
○	Culture change	not in roadmap

Decisions used to require a quarterly re-litigation

BEFORE AND AFTER THE FRAMEWORK



The architecture change is the calibration loop. The culture change is still in Phase 02.

The same transition for the board deck

BEFORE AND AFTER · BANNER-TAG MODIFIER

BEFORE

Every prioritization debate from first principles. Average close: 4 hours, p99 an entire offsite. No audit trail, no calibration — decisions were made, outcomes happened, nobody connected them. We called it "moving fast."



AFTER

Decisions log their rationale; scoring is weighted and calibrated; an audit trail exists. It has 18 entries. We are moving slower now. We are calling this "moving thoughtfully."

The same build decision in banner-tag chrome

DECISION · BANNER-TAG MODIFIER

BUILD

Owns the scoring policy, the calibration loop, and every responsibility the evaluation said "a vendor would own" in the path we rejected.

WHY NOT BUY

Three vendors, none exposing calibration weights — the criterion set by the evaluation team. The evaluation team builds frameworks.

WHY NOT DELAY

The window closes in 18 months. We are in month 7. Phase 02 is planned for month 6. We are aware of this.

Recalibration used to require a coordinated freeze

BEFORE – MANUAL RECALIBRATION

Operators schedule a window, freeze new decisions, swap the weights, verify, then lift the freeze. Average review pause: 18 working hours. Also known as the way we did it last year, which everyone agreed was fine until this deck was written.

AFTER – VERSION-FLOOR RECALIBRATION

The loop emits a new policy with an incremented version; teams pick it up on next refresh. No freeze, no coordinated cutover. This is the good path. It is also the one that has not shipped yet.



How to roll this out across your organization

STEP 01

Pick one team and one decision type

Start with a team that has a real rhythm. No rhythm, no framework.



STEP 02

Log everything, decide nothing differently

Just log, for a month. "Low-effort." Highest dropout rate.



STEP 03

Run your first retrospective

Day 30, score against outcomes. No outcomes? Score the room, call it "qualitative."



STEP 04

Expand to a second team

Your evidence: one team ran one retrospective. Use it confidently.

The weekly practice in three moves

STEP 01

Sense

Observed inputs, never invented. Skipped about 70% of the time.



STEP 02

Score

A signal is data once it carries a number. Calibrating it takes 14 months.



STEP 03

Decide

A signal plus a deadline. The loop closes if anyone logged a prediction.

The case for the framework in three moves

Claim

Calibrated prioritization with audit-grade decision custody. We stop paying the re-litigation cost on every quarterly review.

Evidence

Six months across four product teams. Close-time dropped; calibration ran once. The four teams have since been reorganized into three. None of them count in the "before" baseline.

Implication

The framework works. The deck says so. The deck was written by the framework team. We trust the framework team — the framework told us to.

Wiring a signal into the framework is three lines of application code

JAVASCRIPT • DECISIONFRAMEWORK SDK V2 • 847 TRANSITIVE DEPENDENCIES

```
import { DecisionFramework } from "@company/signal-sdk";

const framework = new DecisionFramework({ configFile: "./framework.config.json" });

// Score a signal at intake
const score = await framework.score(signal, { dimensions: ["confidence", "recency", "relevance"] });

// Log every decision – calibration depends on it
// (nobody calls this line in production, but it is here)
const entry = await framework.decisions.log(decision, { signals: [signal.id], rationale });
```


BEFORE & AFTER • SCORING MECHANICS

Spreadsheet-driven scoring versus framework-driven scoring that is basically also a spreadsheet

BEFORE • THE HONEST SPREADSHEET

```
# Manual scoring. Auditable in the  
# sense that you can see who last  
# edited the file  
import pandas as pd  
  
signals = pd.read_csv('./signals.csv')  
signals['score'] = signals.apply(  
    lambda r: 0.33*r.confidence + 0.33*r.recency + 0.33*r.relevance,  
    axis=1,  
)  
signals.to_csv('./scored.csv')
```

AFTER • THE FRAMEWORK

```
# Calibrated weights, signed policy,  
# every score is audit-logged,  
# same math as the spreadsheet  
from decision_framework import Calibrator  
  
calibrator = Calibrator.load('./policy.json')  
for signal in calibrator.intake.unscored():  
    calibrator.score(signal)  
    calibrator.decisions.log_if_relevant(signal)
```


SECTION 04 PRESENTS SIX MONTHS OF PRODUCTION

DATA. THE DENOMINATOR IS CAREFULLY CHOSEN.

SECTION 03 OF 05 COMPLETE

The framework is built.
Twelve percent of
eligible PMs use it

SECTION 04 · THE RESULTS

Six months of data.
Eighteen logged decisions.
One calibration cycle

The pilot team measured what the pilot team built.

Before the results, what you should have learned in the first 46 slides

- 01 Four components. Two used regularly. → Section 01
- 02 The evaluation team recommended the tool the evaluation team built. → Section 02
- 03 Phase 01 shipped. Phase 03 in the roadmap since 2024. → Section 03
- 04 Adoption is 12%. Target 90%. The gap is "cultural." → this section
- 05 The calibration loop has run once. We call it "calibrated." → this section

Where we are against quarter targets

94%

Signal-classification success

TARGET 99%, GAP IS "KNOWN ISSUE"



18 min

p99 decision close

TARGET 20 MIN, BEATING TARGET

18

Decisions logged

TARGET 340, GAP IS "CULTURAL"


1

Calibration cycles run

TARGET 6, GAP IS "STRUCTURAL"

Phase 01 readiness, by workstream


The status pills reflect the most optimistic reading of the data.



Source: Linear · "blocked" means blocked, we are working on the wording


The same data, dark canvas

Same bars, same numbers. Dark makes the 12% adoption bar feel intentional, not alarming.
That is not why we have a dark modifier. It is why this slide uses it.



Same data, minimal treatment

The lucent strip is gone, so the adoption number is louder. An argument for the non-minimal treatment.




Source: Linear · refreshed 2026-05-07

H1 2026 · 1,840 PERSON-HOURS

Where the engineering quarter went

The "Toil and on-call" wedge is the one nobody put in the roadmap.



Last updated 2026-05-07 · the 9% is probably higher

FRAMEWORK ARC

How the framework arrived in production

Four stages over eighteen months. The connective tissue is described as "momentum."



Cross-functional sign-off · "cross-functional" means two teams instead of one

FINDING 01 • STRUCTURED INTAKE

Structured intake performed above expectations — volume and latency were not the problems we thought

What worked

API connectors handled 94% of structured signals with no manual touch, at 4-minute average latency. This is the part of the demo we show everyone.

What required tuning

NPS verbatim classification ran an 18% error rate for the first two weeks. It's mentioned in the appendix, not in the headline numbers.

KEY INSIGHT

Viable as designed. The headline numbers are accurate. The denominator is carefully chosen.

Three scoring failure modes found in the pilot

- 1

Recency dominance

Recency set above 50% — last week felt more real than six months of NPS. Capped at 40%, until a VP overrides it.
- 2

Source concentration

One customer was 34% of intake. Also 31% of revenue, so the diversity floor waited for renewal.
- 3

Outcome misclassification

"Improve retention" won't score. "Make customers happier" was submitted twelve times, scored 5 each.

WALK THESE QUESTIONS WITH ME IN 60-90 MINUTES.

THE OUTPUT IS A DESIGN WE CAN EXECUTE — OR

AGREEMENT THAT WE NEED ANOTHER SESSION TO

DESIGN IT.

WHAT WOULD HELP US MOVE FORWARD

Next step is a working
session, not a debate

SECTION 05 • THE PATTERN LIBRARY

The story closed on slide 57. These slides demonstrate authoring patterns

The reference appendix: the layout variants and
modifiers used above. Its counter runs
independent of the dark divider's — so it stamps
01, not 05.

DARK VARIANT · CONTENT

The token system handles dark without per-element overrides

Every colour is a custom property that remaps on *dark* — cards, text, borders, the suppressed spectrum bar. Popular with executives who want "more premium." Same content, darker box.

HEADER AND FOOTER

Header stays uppercase — footer renders as written

Set `header:` and `footer:` in frontmatter. The header always uppercases; the footer renders as written — which is how "FINAL v3 USE THIS ONE" reaches the board.

The card stack renders cleanly on dark backgrounds

Every card fills from `--bg-alt`, and its border remaps too — both shift automatically on dark. Most decks don't build this way. Theirs break.

The accent edge keeps `--accent` unchanged, because it already reads on dark. That's partly why the accent was chosen. You're welcome.

Body text becomes a warm light tone via `--text-body` — never pure white. Pure white on dark is technically readable and visually unkind.

Two-card layouts work equally well inverted to dark

The framework asks one calibration question: what proves your scoring model is improving, and what does a quarter without that proof cost? Every other question in this deck depends on the answer. In practice the answer is "we'll revisit it in Q3" — which is also the quarter the workshop to define "improving" is scheduled for.

It's the same shape as any page of written argument — claim, then support. The dark palette changes none of the reading rhythm or the information density. What it changes is how the slide looks in the screenshot your VP forwards to their VP with "see, this is what I mean." The content is identical to the light version. The forward is the feature.

LAYOUT · IMAGE

Image right is the default — text leads, evidence follows

For slides where the visual supports the argument.
Native aspect ratio, no cropping — usually a stock lighthouse called "a metaphor for signal."






MODIFIER · IMAGE MIRROR

Mirror flips the image to the other half — alias of legacy `image left`

`mirror` is the canonical orientation flag. The `left` alias survives for the 40 decks nobody wants to find.



Signal Pipeline · Reference Visualization

Weekly Signal Brief — shipped every Monday, opened by an estimated 3 of 14 eligible PMs

Code chips render inside card titles and body alike

Signal Intake

v2.4

94% of structured signals. The other 6% are CEO Slack messages, escalated unscored.

Scoring Model

configurable

Defaults `33 / 33 / 33`. Most set recency to `99` after a bad quarter, call it calibrated.

Decision Log

required

Every change above `P2` needs a rationale. Introduced Q2 2024. Not yet enforced.

Calibration Loop

monthly

First real update after `2 cycles` — about `14 months`.

A trailing italic-only paragraph becomes an annotation

Signal Intake

Weekly collection across conversations and market data.

Scoring Model

Three dimensions, configurable by seven admins.

Decision Log

Every decision recorded. Enforced by honor system.

Calibration Loop

Predictions vs outcomes. Averages 2.3 attendees.

KEY INSIGHT

The calibration loop separates teams that learn from teams that repeat the same mistakes more thoroughly.

♦ *Pilot retrospective: six months, four teams, one since reorganized out of existence.*

A trailing blockquote becomes a key insight; a plain paragraph becomes a below-note

Signal Intake

The one everyone demos.

Scoring Model

The one everyone reconfigures.

Decision Log

The one nobody fills out.

Calibration Loop

The one that finds no predictions.

KEY INSIGHT

The calibration loop separates teams that learn from teams with very thorough meeting notes about not learning.

Trailing blockquote → key insight; trailing paragraph → below-note.

Compact tightens the spacing scale ~25%

What changes

`--sp-xs` through `--sp-2xl` shrink. Card gaps, gutters, and padding follow, via `var()`.

What does not change

Type, palette, and chrome reservation are untouched. A density flag, not a new layout.

When to reach for it

One card too many, prose a line long, or a denser rhythm.

Composition

Composes with `dark`, `accent`. Not with `title`, `divider`, `image full`.

MODIFIER • LOOSE

Loose is the inverse — more breathing room, same layout machinery

The spacing scale grows ~25%. Reach for *loose* when a slide carries one editorial point and wants to feel quiet.

KEY INSIGHT

Density is not importance. **loose** says this page deserves room — because it carries one thing well. The framework team applied it to this quote, which the framework team wrote about the framework.

Routing signals by source type versus routing by confidence tier

ROUTE BY SOURCE TYPE

Structured signals — NPS, tickets, call notes — go to Intake. Unstructured ones — field notes, conference talk, a board member with dinner instincts — route to the owner for manual classification. Clean in theory, fragile at every edge involving an enterprise customer.



ROUTE BY CONFIDENCE TIER

Signals above 3.0 confidence route straight to scoring; everything below queues for human review. Adaptive — once you've calibrated confidence, which takes two completed cycles. We're in month three, calling it cycle one.

Both paths land on the same PM, who asks why the framework keeps routing things to them.

Chosen flags the right-hand card as the winner

QUARTERLY RECALIBRATION

Weights reviewed once a quarter. Simple, auditable, and blind to everything that happened in between. This is what we did before the framework. We do not discuss that here.



CONTINUOUS CALIBRATION

Weights update at every retrospective once the sample passes 12 logged decisions. We have 18, across six months. The loop has run once.

The right card carries the accent edge and tint — the chosen path. Chosen after the choice was made.

Mirror composes with chosen — the accented card reads from the left

RETROSPECTIVE-TRIGGERED RECALIBRATION

Weights update once the outcome sample passes 12 decisions. `mirror` puts this card on the left; `chosen` accents it. The audience reads the conclusion first — useful when you're confident they won't ask how many decisions are in the sample. There are 18.



QUARTERLY RECALIBRATION

Reviewed once a quarter on a fixed schedule, no outcome data required to trigger it. Source order is preserved — the first card is always the considered alternative. Mirror flips the rendering, not the intent.

The below-note holds the caveat. The loop ran once. We call it continuous.

Mirror puts the hero card on the right; sub-cards stack on the left

The scoring model requires the loop

No calibration, static weights — a spreadsheet. Sub-cards go left, written to make the hero inevitable.


The Decision Log enables the loop

Loop needs outcomes; outcomes need logged decisions. 18 entries, one run. The card the presenter skips.

The calibration loop is what closes the framework

The only component that makes the model improve instead of documenting failures better. `mirror` puts it right — premise before conclusion.

Four switches to 4 columns; pair with compact for visual balance



Glossary

A – F

TERM	DEFINITION
Adoption	PMs logging within 24h of a close. 5.3% now; target 90%. The gap is "cultural."
Auditability	Reconstruct any decision months later without its author. The author is always there anyway.
Calibration	Predicted vs observed. Requires predictions. See "Decision Log," "Adoption," "Gap."
Connector	Sprig ↔ your NPS / support platforms. In "final testing" since Q4 2024.
Decision Log	The append-only record of every decision. Append-only in theory, empty in practice.
Eligible PM	Past 30-day onboarding on an adopted team. 14 exist; 13 think they're still onboarding.
Framework	Speed, Auditability, Adoption, Calibration. Also what people call a spreadsheet.

Glossary

P – T

TERM	DEFINITION
Predicted outcome	The expectation logged at decision time, fed to calibration. Hypothetically.
Prioritization rhythm	Documented "weekly." Observed "whenever there's an incident or a board meeting."
Retrospective	Day-30 scoring of decisions vs outcomes. Often day 45. Always runs long.
Signal	A survey, an NPS comment, a ticket. Or a 9 PM Slack from whoever last saw a customer.
Sprig	The micro-survey behind Tool D. See "Conflict of Interest" — not in this glossary.
Tool D	The recommended option, built by this evaluation's authors, best by their criteria.

MODIFIER · ACCENT

Accent swaps the rainbow stripe for one editorial colour

It tints the heading and composes with `dark`, where `accent.dark` restores the suppressed stripe. The framework team is very thorough.