

DECISION FRAMEWORK • Q3 2025

From Signal to Strategy

A 78-slide answer to the question "have you considered writing things down"

What this deck covers, in order

01 Why This Exists — slide 3

→ **The Framework — slide 7**

03 The Evaluation — slide 22

04 The Build — slide 31

05 The Results — slide 47

CONTEXT • COMPETITIVE DYNAMICS

The window for differentiation is narrowing

Three converging forces — commoditized infrastructure, compressed release cycles, and rising customer switching costs — have reduced the average durable advantage window from 36 months to under 14. Teams that cannot identify signal from noise in that window will consistently miss timing. This slide will appear in every deck for the next two years regardless of whether the window keeps narrowing.

“

The signal was always there. We just didn't have a system that forced us to look at it before we'd already decided.

”

— Head of Product, Pilot Team 3, in a retrospective where we then decided what we had already decided

IMPACT · PILOT RESULTS

Six months of results across four product teams

MEASURED AGAINST PRE-FRAMEWORK BASELINE, SAME TEAMS, SAME CONDITIONS, SAME SPREADSHEET.

73%

FASTER DECISION CLOSE

4.2×

SIGNAL RECALL

18

DECISIONS LOGGED

91%

TEAM ALIGNMENT

CALIBRATION RESULT • 6-MONTH PILOT

14X

Return on signal investment —
calculated by the team that built the
framework, using a baseline they
defined, in a pilot they designed, to
measure an outcome they predicted.
Independently verified by nobody.

We built a four-
component scoring
system. All four
components are
described in the next
fourteen slides. Two of
them are used
regularly


SIGNAL DEFINITION • WORKSHOP 04

Before we score signals, we need to agree on what a signal is

This has taken three workshops. We are now in workshop four. The output of this workshop is a shared definition that will be socialized in a fifth workshop before it can be used.

How signals move from input to decision

Four-stage processing pipeline — 11-week implementation, still in pilot



The framework has four components

Signal Intake

Weekly structured collection across customer conversations, market data, and competitive moves. Normalized into a common schema. This is the part everyone agrees is a good idea but nobody does on the week of the retrospective.

Scoring Model

Each signal scored on three dimensions: confidence, recency, and strategic relevance. Weights are team-configurable, which means the head of product reconfigures them after every retrospective until the output agrees with their roadmap.

Decision Log

Every decision recorded with its signals, options, and criteria applied. Required artifact. Currently has 18 entries. Approximately 340 decisions have been made.

Calibration Loop

Monthly retrospective that compares predicted outcomes to actual outcomes. The meeting exists. Predicted outcomes are rarely logged before the meeting.

Signal Intake produces three outputs

1

Weekly Signal Brief

A ranked list of the top 10 signals from the prior week, with confidence scores and source attribution.

Distributed to product leads every Monday morning, where it sits unread in an inbox folder called "Framework Stuff."

2

Anomaly Alerts

Real-time flags when a signal exceeds the 2σ threshold on any dimension. Routed directly to the accountable PM with a 4-hour response SLA. The PM usually responds in 4 hours to ask what 2σ means.

3

Monthly Signal Index

The source of truth for the calibration loop. A complete record of all signals logged, scored, and resolved in the prior month. Required reading before each retrospective. Nobody has read it. It is comprehensive.

The three things the framework connects

Signal

The observed input. A verbatim, a metric move, a competitor announcement. The unit of intake. Frequently confused with "things the VP heard at a conference."

Decision

A signal plus a deadline. Logged with rationale and predicted outcome. In the framework, every decision is traceable to its signals. In practice, the signal is often "we discussed it at the offsite."

Outcome

The observed result, compared to the predicted outcome at retrospective. The unit of calibration. Currently 18 outcomes have been logged. Approximately 340 outcomes have occurred.

Two failure modes the framework is designed to prevent

False signal amplification.

A single loud voice — one enterprise customer, one analyst report, one VP who "has a feeling" — dominates the decision without being weighed against the full signal set. The scoring model prevents any single source from exceeding 30% of the total signal weight. Unless that source is the CEO, in which case the weight cap is a guideline.

Signal hoarding.

Teams collect signals but do not log decisions, so the calibration loop has nothing to learn from. The Decision Log is a required artifact for any prioritization change above P2 severity. No log, no change. This rule was printed on a poster and placed in the meeting room. The poster has been replaced by a free-pizza flyer.



SECTION 01 • CONTINUED

Scoring Model Deep Dive

Dark panel + content • split-panel watermark

What the scoring model actually does

The scoring model is the most configurable component, which is either a feature or a warning sign depending on whether you trust your team to configure things without turning every dimension to 100. This section covers the three dimensions.

1

Confidence

How many independent sources corroborate the signal. Ranges 1–5. Enterprise customers count as 1, regardless of volume.

2

Recency

Time-decay from signal date to scoring date. Half-life is team-configurable. Most teams set it to two weeks and then wonder why they only act on recent news.

3

Strategic Relevance

Manual score from the signal owner. Ranges 1–5. Requires written justification above 4. Scores of 5 have a remarkable correlation with whoever is presenting the roadmap this quarter.

The six signal dimensions, what they measure, and how they are scored

01	Confidence	Number of independent sources corroborating the signal	1-5 · AUTO-SCORED · ENTERPRISE CUSTOMERS ALWAYS GET A 4
02	Recency	Time-decay from signal date, configurable half-life	0.0-1.0 · AUTO-SCORED · TEAMS SET HALF-LIFE TO "VERY SHORT" AFTER A BAD QUARTER
03	Relevance	Alignment to current strategic bets, owner-scored	1-5 · MANUAL · CORRELATES STRONGLY WITH WHAT THE PM ALREADY PLANNED TO BUILD
04	Reach	Number of customers or segments affected	1-5 · AUTO-SCORED · USUALLY 5 IF THE ENTERPRISE CUSTOMER MENTIONED IT
05	Effort	Engineering and design cost to act on the signal	1-5 · MANUAL · SCORED BY PMS · REVIEWED WITH RAISED EYEBROWS BY ENGINEERING
06	Confidence delta	Change in confidence score since last scoring cycle	-5 TO +5 · AUTO · MOSTLY MEASURES WHETHER SOMEONE TALKED TO A CUSTOMER THIS WEEK

Scoring model: before and after the calibration loop

BEFORE CALIBRATION

Equal weights. Confidence, recency, and relevance each contribute 33%. Simple and honest about the fact that we are basically guessing.



AFTER CALIBRATION

Weights reflect your team's historical signal accuracy, except that the team keeps changing and the history is three months of data collected during a quarter that everyone agrees was atypical.

The shift from equal weights to calibrated weights takes two retrospective cycles — roughly 60 days from adoption, or 14 months, depending on who you ask.

Two intake modes for different signal types

Structured Intake

Signals with clear schema: NPS verbatims, support tickets, feature request volumes, win/loss notes.

Ingested automatically. Scored on arrival. Zero manual handling. Produces 94% of the data and 12% of the roadmap decisions.

Unstructured Intake

Signals without schema: field observations, conference conversations, analyst briefings, a board member who attended a dinner and has thoughts. Require human classification. Routed to the signal owner for a 48-hour classification window. Produce 6% of the data and 88% of the roadmap decisions.

How a decision moves through the framework



What the framework does not do

It does not make decisions — it structures the information that humans use to make the decisions they were going to make anyway.

It does not replace customer discovery — it scores and routes what discovery surfaces, if discovery surfaces anything before the roadmap is already locked.

It does not work without the Decision Log — calibration requires outcome data. The Decision Log currently has 18 entries. We have made approximately 340 decisions.

It does not guarantee alignment — it surfaces disagreement earlier, so you can have the same argument at the beginning of the quarter instead of the end.

It does not scale down to individual feature decisions — it is designed for prioritization above P2, so the 90% of daily decisions happening below that threshold are still being made in Slack threads from 2022.

Four things that must be true before you begin

- 01 You have a regular prioritization cadence — at minimum monthly. If your cadence is "whenever someone escalates loudly enough," the framework will make that escalation louder.
- 02 At least one person owns signal collection full-time or as a primary responsibility. "Primary" means "more than 10% of their time," which means you need to write it into their job description, which means you need to agree on what a signal is first.
- 03 Leadership has agreed to log decisions with rationale, not just outcomes. This agreement is harder than it sounds. This slide exists because it is harder than it sounds.
- 04 You have 90 minutes per week. Budget four.

SECTION 02 EVALUATES FOUR TOOLS AGAINST
FOUR CRITERIA. THE EVALUATION TEAM BUILT
ONE OF THE TOOLS. THE CRITERIA WERE
DEFINED BY THE EVALUATION TEAM. THIS IS
DISCLOSED IN A FOOTNOTE.

SECTION 01 OF 05 COMPLETE

The framework is
specified. The next
question is whether
to buy one instead

SECTION 02 · THE EVALUATION

We evaluated four
tools. One of them was
built by the evaluation
team

Four requirements every decision system must meet

01

Speed

Decisions must close within the window they are relevant to. A framework that takes six months to calibrate before its first useful output is not fast. We are not going to address this directly on this slide.

02

Auditability

Every prioritization decision above a threshold must carry a traceable rationale. Required for alignment, compliance, and for reconstructing what happened after the launch goes badly.

03

Adoption

If the team won't use it weekly, calibration never runs. Ninety minutes per PM is the ceiling we put in the deck. The actual usage data suggests the ceiling is closer to 11 minutes and one skipped retrospective per month.

04

Calibration

The system must improve over time. A static scoring model is a spreadsheet with extra steps and a dashboard nobody checks.

We evaluated four intake tools against the criteria

Tool A · Chorus

- ✓Speed
- ✖Auditability
- ✓Adoption
- ✖Calibration

Strong call recording. No decision logging. No calibration.
Recommended by the sales team because they already use it and didn't want another tool.

Tool B · Productboard

- ✖Speed
- ✓Auditability
- ✓Adoption
- ✖Calibration

Solid intake and prioritization. No calibration mechanism. Setup takes 3–4 weeks, which we said was fine and then spent 11 weeks on.

Tool C · Notion

- ✓Speed
- ✓Auditability
- ✖Adoption
- ✖Calibration

Flexible enough to build anything. Teams built seven slightly different versions of the framework and then debated which one was canonical for two quarters.

Tool D · Sprig + Decision Log

- ✓Speed
- ✓Auditability
- ✓Adoption
- ✓Calibration

Built by the team that wrote this evaluation. Meets all four criteria as defined by the team that built it. Recommended.



The four tools side by side

CRITERION	CHORUS	PRODUCTBOARD	NOTION	SPRIG + LOG
Speed	✓	✗	✓	✓
Auditability	✗	✓	✓	✓
Adoption	✓	✓	✗	✓
Calibration	✗	✗	✗	✓
Setup time	1 day	3–4 weeks	40+ hours	Same day

♦ Evaluated against criteria defined by the team building Sprig + Log. We want to be transparent about this. It is in the footnote.

How we sort the four tools against our two axes

COVERAGE · COST

High coverage / Low cost

Sprig + Log — strongest fit on coverage, lowest TCO.

Recommended. Built by the evaluation team.

Productboard — narrower coverage, mid-tier license, also on this list to show we looked.

High coverage / High cost

Notion build-out — full coverage theoretically. Premium maintenance cost. We have tried this. The result is seven slightly different versions of the framework and a quarterly debate about which is canonical.

Low coverage / Low cost

Chorus — cheap, leaves three criteria uncovered, popular anyway because the sales team already uses it and has strong feelings about switching costs.

Low coverage / High cost

*None — and that is the signal. Or a gap in our vendor universe.
We are treating it as a signal.*

Applying the criteria to the tools — here is where the evidence points

The evidence favors Tool D



Sprig combined with a lightweight Decision Log meets all four criteria within the 90-minute weekly budget, reaches production in the same week it is adopted, and was recommended before the evaluation began. The evaluation confirmed this.

The path is not self-executing

Sprig requires a connector to your NPS and support platforms. Budget 4–6 hours of engineering time in week one. The connector has been in review for 11 weeks. We are still calling this "week one."

The Decision Log is the hardest part

Not technically. Culturally. PMs need to log decisions before they close, not after. This is a habit change. It has been described as a habit change in every deck since Q3 2023. It remains a habit change.

The evaluation came down to one question the vendors could not answer

BUY A VENDOR FRAMEWORK

Three vendors evaluated. None expose calibration weights to the customer, which is the criterion that eliminated all three. This criterion was added to the evaluation after the evaluation team decided to build.



BUILD THE FRAMEWORK IN-HOUSE

Owens the scoring policy, owns the calibration loop, owns the timeline. The competitive window closes in 18 months. A vendor cutover would consume nine of those, which is how we knew building was faster before the evaluation began.

The left card is struck through. The DECISION connector is bold. The conclusion was reached before the slide was built. The slide was built so the conclusion would have a slide.

We are building, not buying

DECISION · 2026 Q1

BUILD

Owens the scoring policy, owns the calibration loop, owns the timeline. Also owns the maintenance, the on-call rotation, and any future explanation of why the framework scored the wrong thing.

WHY NOT BUY

Three vendors evaluated. None expose calibration weights to the customer. The weights are the product. We cannot buy the product without the weights. This was the finding.

WHY NOT DELAY

The competitive window closes in 18 months. This sentence has been in the deck since Q1 2025.

How we make calls when the spec is silent

- 01 We default to the choice that is cheaper to reverse, unless reversing it would require a meeting, in which case we pick something and document it as a decision.
- 02 We name the actor, never the system. The system cannot be held accountable. The actor can be reorganized.
- 03 We write down the bet on the same slide as the choice. We do not always write down the slide number where the bet will be reviewed.

What the build covers

Three phases, four workstreams. We own the policy, the loop, and the timeline. We also own whatever Phase 2 turns into.

1

Phase 01 — Architecture

Scoring policy live. Decision Log accepting entries. One pilot team in weekly cadence.

2

Phase 02 — Operations

Multi-team calibration. Automated weight updates. Before-after data sufficient to defend in a board update.

3

Phase 03 — Scale

Org-wide enablement. This phase has been in the roadmap since 2024. We remain committed.

P

SECTION 03 • THE BUILD

Phase 01 through
Phase 03 ships the
architecture, the
operations, and the
scale

What ships in each phase, by workstream

WORKSTREAM	PHASE 01	PHASE 02	PHASE 03
Signal Intake	Connector v1	Multi-source dedupe	Anomaly auto-routing
Scoring	Equal-weights model	Per-team calibration	Per-decision profiles
Decision Log	Append-only schema	Outcome auto-pairing	Examiner export
Adoption	One pilot team	—	Org-wide enablement

Phase 3 is where org-wide enablement lives. Phase 3 has been in the roadmap since 2024. Phase 2 is ongoing.

Who owns each part of the framework lifecycle

Signal custody

Signal owner

Manages intake quality and source diversity. Never tunes scoring weights directly. Tunes them indirectly by choosing which signals to surface and which to classify as low-confidence.

Policy

Framework operator

Owns scoring policy, calibration cadence, version floors, and rollback playbooks. Currently one person. This is noted in the risk register.

Consumption

Product team

Holds time-bound scoring profiles. Runs intake and decision-logging. Primarily interacts with the framework by asking the Framework operator to adjust the weights.

Oversight

Auditor

Reads the Decision Log audit trail. Cannot edit weights. Has read the audit trail once. Found 18 entries. Asked if that was expected.

Three phases get us from decision to org-wide adoption

PHASE 01

Architecture

The first phase scopes what we build, what we buy, and what we defer. Output is an architecture decision record. The decision record will itself require an architecture phase before it can be approved.

PHASE 02

Pilot

One team, one decision type, one quarter. The phase ends when the framework is in production cadence. "Production cadence" means the retrospective happened at least once.

PHASE 03

Rollout

Five teams in two months. The phase ends when adoption is above 90% and no team needs handholding. This phase is currently planned for Q2. It has been planned for Q2 for three consecutive years.

Three milestones mark Phase 01 complete

MILESTONE A

Scoring policy in production

The signed scoring policy runs end-to-end. The first calibrated brief lands in leadership's inbox. Leadership replies with a question about whether the weights can be adjusted. They can.

MILESTONE B

Per-team weights

One framework can carry distinct scoring weights per team without per-team forks. Recalibration is a single policy update. Each team will still want their own.

MILESTONE C

Per-decision-class profiles

Authoring a scoring profile scoped to a single decision class takes minutes. Agreeing on what constitutes a decision class takes a workshop series. See slide 8.

PHASE 01 · ACCEPTANCE REVIEW

What shipped, what slipped, what stayed open

- ☒ Scoring policy live across all pilot teams
- ☒ Decision Log audit trail readable by Auditor role shipped 2026-Q1
- ☒ One reference team running end-to-end weekly cadence one team
- ☐ Examiner pack auto-generation was Phase 01 · now Phase 1b
- ☐ Multi-team calibration cadence Phase 02
- ☐ Adoption above 90% Phase 03

- ☐ Culture change not in roadmap

Decisions used to require a quarterly re-litigation

BEFORE AND AFTER THE FRAMEWORK

BEFORE

Every prioritization debate started from first principles. Average close: 4 hours. p99: an entire offsite. Outcome: whatever the most senior person in the room wanted, expressed as a consensus.



AFTER

Decisions resolve against logged weights and prior calibration outcomes. Average close: 18 minutes. Outcome: whatever the scoring model outputs, after the weights are adjusted to reflect what the most senior person in the room wants.

The architecture change is the calibration loop. The culture change is still in Phase 02.

The same transition for the board deck

BEFORE AND AFTER · BANNER-TAG MODIFIER

BEFORE

Every prioritization debate from first principles.
Average close: 4 hours. No audit trail. No calibration.
Decisions were made, outcomes happened, nobody connected them. This was called "moving fast."



AFTER

Decisions log their rationale. Scoring is weighted and calibrated. An audit trail exists. The audit trail has 18 entries. We are moving slower. We are calling this "moving thoughtfully."

The same build decision in banner-tag chrome

DECISION · BANNER-TAG MODIFIER

BUILD

Owns the scoring policy, the calibration loop, and all the things the evaluation said "a vendor would own" in the rejected path.

WHY NOT BUY

Three vendors evaluated. None expose calibration weights. The evaluation team defined "exposes calibration weights" as a required criterion. The evaluation team builds frameworks.

WHY NOT DELAY

The window closes in 18 months. We are in month 7. There are 11 months left. Phase 02 is planned for month 6. We are aware of this.

Recalibration used to require a coordinated freeze

BEFORE — MANUAL RECALIBRATION

Operators schedule a recalibration window, freeze new decisions, swap weights, verify, lift the freeze. Average review pause: 18 working hours. Also known as: the way we did it last year, which everyone agreed was fine until this deck was written.



AFTER — VERSION-FLOOR RECALIBRATION

The calibration loop emits a new policy with an incremented version. Teams pick up the new policy on next refresh. No freeze. No coordinated cutover. This is the good path. It is also the one that has not shipped yet.

How to roll this out across your organization

STEP 01

Pick one team and one decision type

Start with a team that already has a regular prioritization rhythm. If no team has one, the framework cannot help you. This is step one.

STEP 02

Log everything, decide nothing differently

In the first month, do not change how you make decisions. Just log them. This is described as low-effort. It is the step with the highest dropout rate.

STEP 03

Run your first retrospective

At day 30, score the logged decisions against outcomes. If no outcomes have been logged, score the energy in the room and call it "qualitative calibration."

STEP 04

Expand to a second team

With one retrospective complete, you have evidence. The evidence is that one team ran one retrospective. Use it confidently.

The weekly practice in three moves

STEP 01

Sense

Inputs are signals. Signals are observed, never invented. The first step is to write down what you see, not what you conclude from it. This step is skipped approximately 70% of the time.

STEP 02

Score

A signal becomes data once it carries a number. The number is calibrated against outcomes, not against intuition. The number is set by intuition until enough outcomes exist to calibrate against, which takes two cycles, which takes 14 months.

STEP 03

Decide

A decision is a signal plus a deadline. Without the deadline it is an opinion. The retrospective closes the loop, provided anyone logged a predicted outcome, which brings us back to step one.

The case for the framework in three moves

Claim

The framework gets calibrated prioritization with audit-grade decision custody. We stop paying re-litigation cost on every quarterly review.

Evidence

The pilot ran six months across four product teams. Decision close-time dropped. Calibration ran once. The four teams have since been reorganized into three teams. One is now a platform team. None of them count in the "before" baseline.

Implication

The framework works. The deck says so. The deck was written by the framework team. We trust the framework team. The framework told us to.

Wiring a signal into the framework is three lines of application code

JAVASCRIPT · DECISIONFRAMEWORK SDK V2 · 847 TRANSITIVE DEPENDENCIES

```
import { DecisionFramework } from "@company/signal-sdk";

const framework = new DecisionFramework({ configFile: "./framework.config.json" });

// Score a signal at intake
const score = await framework.score(signal, { dimensions: ["confidence", "recency", "relevance"] });

// Log every decision – calibration depends on it
// (nobody calls this line in production, but it is here)
const entry = await framework.decisions.log(decision, { signals: [signal.id], rationale });
```


BEFORE & AFTER • SCORING MECHANICS

Spreadsheet-driven scoring versus framework-driven scoring that is basically also a spreadsheet

BEFORE • THE HONEST SPREADSHEET

```
# Manual scoring. Auditable in the  
# sense that you can see who last  
# edited the file  
import pandas as pd  
  
signals = pd.read_csv('./signals.csv')  
signals['score'] = signals.apply(  
    lambda r: 0.33*r.confidence + 0.33*r.recency + 0.33*r.relevance,  
    axis=1,  
)  
signals.to_csv('./scored.csv')
```

AFTER • THE FRAMEWORK

```
# Calibrated weights, signed policy,  
# every score is audit-logged,  
# same math as the spreadsheet  
from decision_framework import Calibrator  
  
calibrator = Calibrator.load('./policy.json')  
for signal in calibrator.intake.unscored():  
    calibrator.score(signal)  
    calibrator.decisions.log_if_relevant(signal)
```


SECTION 04 PRESENTS SIX MONTHS OF
PRODUCTION DATA. THE CALIBRATION LOOP HAS
RUN ONCE. THE ADOPTION GAP IS DESCRIBED
AS CULTURAL. THE DATA IS ACCURATE. THE
DENOMINATOR IS CAREFULLY CHOSEN.

SECTION 03 OF 05 COMPLETE

The framework is
built. The scoring
policy is live. Twelve
percent of eligible
PMs are using it

SECTION 04 • THE RESULTS

Six months of data.
Eighteen logged
decisions. One
calibration cycle

The pilot team measured what the pilot team built.

Before the results, what you should have learned in the first 46 slides

- 01 The framework has four components. Two are used regularly. → Section 01
- 02 The evaluation team recommended the tool the evaluation team built. → Section 02
- 03 Phase 01 shipped. Phase 02 is planned for Q2. Phase 03 has been in the roadmap since 2024. → Section 03
- 04 Adoption is at 12%. The target is 90%. The gap is described as cultural. → this section
- 05 The calibration loop has run once. We are calling this "calibrated." → this section

Where we are against quarter targets



94%

Signal-classification success

TARGET 99%, GAP IS "KNOWN ISSUE"

18 min

p99 decision close

TARGET 20 MIN, BEATING TARGET

18

Decisions logged

TARGET 340, GAP IS "CULTURAL"


1

Calibration cycles run

TARGET 6, GAP IS "STRUCTURAL"

Phase 01 readiness, by workstream


Snapshot taken at 14:00 UTC. Status pills reflect the most optimistic interpretation of available data.



Source: Linear · refreshed 2026-05-07 · "blocked" means blocked, we are working on the wording


The same data, dark canvas

The bars, status pills, and numbers are identical. The dark canvas makes the 12% adoption bar feel more intentional and less alarming. This is not why we have a dark modifier. It is why this slide uses it.



Same data, minimal treatment

The lucent strip is gone. The adoption number is more visible. This is an argument for the non-minimal treatment.




Source: Linear · refreshed 2026-05-07

H1 2026 · 1,840 PERSON-HOURS

Where the engineering quarter went

Wedges drawn proportionally. The "Toil and on-call" wedge is the one nobody put in the roadmap.




- Signal Intake build 46%
- Scoring policy work 22%
- Decision Log integration 18%
- Explaining the framework to stakeholders 9%
- Toil and on-call 5%

Refreshed weekly · last updated 2026-05-07 · the 9% is probably higher

How the framework arrived in production

Four stages over eighteen months. Date pill leads each item; status pill trails. The connective tissue between stages is described as "momentum."



Cross-functional sign-off · 2026-04-29 · "cross-functional" means two teams instead of one

FINDING 01 • STRUCTURED INTAKE

Structured intake performed above expectations — volume and latency were not the problems we thought they were

What worked

API connectors handled 94% of structured signals without manual intervention. Average scoring latency was 4 minutes. Schema normalization held. This is the part of the demo we show everyone.

What required tuning

NPS verbatim classification had an 18% error rate in the first two weeks. The classification model needed training. This is mentioned briefly in the appendix and not included in the headline numbers.

KEY INSIGHT

Viable as designed. The headline numbers are accurate. The denominator is carefully chosen.

Three scoring failure modes found in the pilot

1

Recency dominance

High-recency noise crowding out durable signal. Teams set recency weight above 50% in the first calibration pass because last week's customer call felt more real than six months of NPS data. Corrected by capping recency weight at 40% until two calibration cycles complete, or until a VP overrides the cap, whichever comes first.

2

Source concentration

Single-customer signals inflating confidence scores. One enterprise customer's verbatims represented 34% of all structured intake in month one. The customer represented 31% of revenue, so nobody wanted to add the source-diversity floor until after renewal.

3

Outcome misclassification

PMs logging predicted outcomes too vague to score at retrospective. "Improve retention" is not scoreable. "Reduce 30-day churn from 8.2% to below 7%" is scoreable. "Make customers happier" was submitted twelve times and scored 5 on strategic relevance every time.

WALK THESE QUESTIONS WITH ME IN 60-90
MINUTES. THE OUTPUT IS EITHER A DESIGN WE
CAN EXECUTE, OR A SHARED UNDERSTANDING
THAT WE NEED ANOTHER WORKING SESSION TO
DESIGN THE DESIGN. EIGHTEEN DECISIONS
LOGGED. THREE HUNDRED FORTY MADE. WE HAVE
A FRAMEWORK FOR THIS.

WHAT WOULD HELP US MOVE FORWARD

Next step is a
working session, not
a debate

SECTION 05 • THE PATTERN LIBRARY

The story closed on slide 57. These slides demonstrate authoring patterns

The narrative arc ended at the working-session ask. Section 05 is the reference appendix — layout variants and modifiers used throughout the deck, demonstrated for authors who want to replicate them. The divider light counter is independent of the dark-divider counter, which is why this slide stamps 01 rather than 05.

DARK VARIANT • CONTENT

The token system handles dark without per-element overrides

All colours reference CSS custom properties that remap when *dark* is added. Cards, headings, borders, and body text all shift automatically. The spectrum bar is suppressed on dark slides. Dark mode is popular in decks presented to executives who have asked for something that "feels more premium." It is the same content in a darker box.

HEADER AND FOOTER

Header stays uppercase — footer renders as written

Set `header:` and `footer:` in frontmatter for deck-level labels. The header always uppercases automatically, which means you can write it in title case and it will look like you used all-caps intentionally. The footer renders exactly as written, which is why some footers say "FINAL v3 FINAL USE THIS ONE" and nobody notices until the board meeting.

The card stack renders cleanly on dark backgrounds

Every card uses `--bg-alt` for fill and `--border` for the border — both remap in dark mode. This is the correct way to build a system. Most decks do not do this. Theirs break.

The accent left border uses `--accent` which is unchanged — it reads well against dark. The accent colour was chosen partly for this reason. You are welcome.

Body text shifts to `--text-body` which in dark mode is a warm light tone, not pure white. Pure white body text on a dark slide is technically readable and visually unkind.

Two-card layouts work equally well inverted to dark

The framework introduces a single calibration question: what proves your scoring model is improving, and what is the cost of running without that proof for a quarter? Every other question in this document depends on the answer. The answer, in practice, is that we will revisit this in Q3.


The pattern here is the same as any page of written argument — claim, then support. The dark palette does not change the information density or the reading rhythm. It does change how the slide looks in the screenshot your VP sends to their VP with the message "see, this is what I mean."

LAYOUT · IMAGE

Image right is the default — text leads, evidence follows

This is the pattern for slides where the visual is the supporting material, not the argument. The image preserves its native aspect ratio; the lattice pattern frames whatever bands remain. No cropping, ever. Authors see the image they dropped in, which has historically been a stock photo of a lighthouse described as "a metaphor for signal."






MODIFIER · IMAGE MIRROR

Mirror flips the image to the other half — alias of legacy `image left`

The `mirror` modifier is the canonical cross-cutting orientation flag. The `left` alias is preserved for backwards compatibility, which is the polite way of saying there are 40 existing decks that use `image left` and nobody wants to find them.



Signal Pipeline · Reference Visualization

Weekly Signal Brief — the intake pipeline's primary deliverable, shipped every Monday,

11:00 AM PST (7:00 AM GMT) on the 1st of each month

Code chips render inside card titles and body alike

Signal Intake

v2.4

Handles 94% of structured signals without manual intervention. The other 6% are Slack messages from the CEO that get escalated directly to the roadmap with zero scoring.

Scoring Model

configurable

Default weights are 33 / 33 / 33 — adjust after your first retrospective. Most teams set recency to 99 after a bad quarter and call it calibrated.

Decision Log

required

Every prioritization change above P2 must carry a logged rationale. No log, no change. This rule was introduced in Q2 2024. It has not yet been enforced.

Calibration Loop

monthly

First meaningful weight update happens after 2 cycles. The two cycles take approximately 14 months to complete in practice.

A trailing italic-only paragraph becomes an annotation

Signal Intake

Weekly structured collection across customer conversations, market data, and competitive moves.

Scoring Model

Three dimensions: confidence, recency, and strategic relevance. Weights configurable by anyone with admin access, which is currently seven people.

Decision Log

Every decision recorded with its signals and the criteria applied. Required. Enforced by honor system.

Calibration Loop

Monthly retrospective comparing predictions to outcomes. Averages 2.3 attendees.

KEY INSIGHT

The calibration loop is what separates teams that learn from teams that repeat the same mistakes while documenting them more thoroughly.

A trailing blockquote becomes a key insight; a plain paragraph becomes a below-note

Signal Intake

Weekly structured collection across customer conversations, market data, and competitive moves.
The one everyone demos.

Scoring Model

Three dimensions: confidence, recency, relevance. The one everyone reconfigures.

Decision Log

Every decision recorded with its signals and criteria.
The one nobody fills out.

Calibration Loop

Monthly retrospective comparing predictions to outcomes. The one that finds there are no predictions to compare.

KEY INSIGHT

The calibration loop is what separates teams that learn from teams that have very thorough meeting notes about not learning.

Compact tightens the spacing scale ~25%

What changes

`--sp-xs` through `--sp-2xl` shrink. Card gaps, list gutters, and section padding follow because every layout reads them via `var()`.

What does not change

Type ramp, palette, and chrome reservation (header / footer / pagination) are untouched. Compact is a density flag, not a different layout.

When to reach for it

You have one more card than fits, or your prose runs the section by 1–2 lines, or you want a denser visual rhythm without rewriting copy.

Composition

`compact` composes with `dark`, `accent`, and any layout where density makes sense. It is silently incompatible with `title`, `divider`, and `image full`.

Loose is the inverse — more breathing room, same layout machinery

The spacing scale grows ~25% rather than shrinks. Reach for *loose* when the slide carries a single editorial point and you want the page to feel deliberately quiet — values pages, declarative principles, the one thing you actually want the audience to remember.

KEY INSIGHT

Density is not the same as importance. **loose** says: this page deserves room — not because it carries more, but because it carries one thing well. The framework team has applied **loose** to this quote. The quote was written by the framework team about the framework.

Routing signals by source type versus routing by confidence tier

ROUTE BY SOURCE TYPE

Structured signals — NPS, tickets, call notes — go to Intake. Unstructured signals — field observations, conference conversations, a board member who attended a dinner and has strategic instincts — go to the signal owner for manual classification. Clean conceptually. Fragile at every edge case involving an enterprise customer.



ROUTE BY CONFIDENCE TIER

Signals above 3.0 confidence route directly to the scoring model. Below-threshold signals queue for human review. Adaptive, provided you have calibrated confidence scoring first. Calibrated confidence scoring requires two completed retrospective cycles. We are in month three. We are calling it cycle one.

The below-note is for the finding that did not fit the cards. Both routing paths eventually land on the same PM who asks why the framework keeps routing things to them.

Chosen flags the right-hand card as the winner

QUARTERLY RECALIBRATION

Weights are reviewed once a quarter. Simple.
Auditable. Blind to everything that happened in
between. This is what we did before the framework.
We do not discuss that in this deck.



CONTINUOUS CALIBRATION

Weights update at every retrospective once a
minimum sample is reached. The minimum sample is
12 logged decisions. We have 18 logged decisions
across six months. The calibration loop has run once.

The right card carries an accent left-edge and accent-tinted background. It is the chosen path. The deck was written after the choice was made.

Mirror composes with chosen — the accented card reads from the left

QUARTERLY RECALIBRATION

Weights reviewed once a quarter on a fixed schedule. No outcome data required to trigger the review. Simple. Auditable. Source order preserved: the first card is always the considered alternative. Mirror flips rendering, not intent. This is what we did before the framework, which is not stated explicitly on this slide.



RETROSPECTIVE-TRIGGERED RECALIBRATION

Weights update when the outcome sample passes the 12-decision threshold. `mirror` puts this card on the left; `chosen` accents it. The audience reads the conclusion first, which is useful when you are confident they will not ask how many decisions are in the sample. There are 18.

The below-note is for the caveat. The calibration loop has run once. We are calling this continuous.

Mirror puts the hero card on the right; sub-cards stack on the left

The scoring model requires the loop

Without calibration, the weights are static. Static weights are a spreadsheet. The featured layout puts supporting cards on the left. They are written after the hero claim is decided, to make the claim feel inevitable.

The Decision Log enables the loop

The loop requires outcome data. Outcome data requires logged decisions. The Decision Log has 18 entries; the loop has run once. The framework is calibrated. The third sub-card is the one the presenter skips when running long — it holds the most important

The calibration loop is what closes the framework

The loop is the only component that makes the scoring model improve rather than document the same errors with better formatting. Mirror puts the hero card on the right. Use it when the audience needs the premise before the conclusion will land — which is when the premise is "the loop requires logged outcomes."

Four switches to 4 columns; pair with compact for visual balance

Sense

Observe. Write it down.
Do not interpret yet. This
step is the hardest.

Score

A signal is data once it
has a number. The
number is calibrated.
Eventually.

Decide

A decision is a signal plus
a deadline. A recurring
meeting is not a deadline.

Review

The retrospective. The
loop closes here. Or it
would, if the loop had
data.

Glossary

A – F

TERM	DEFINITION
Adoption	Percentage of eligible PMs filing a Decision Log entry within 24 hours of a decision close. Currently 5.3%. Target is 90%. Gap is "a cultural issue."
Auditability	The property that any decision can be reconstructed from its inputs three months later without the original author present. In practice, the original author is always present and explains what the inputs meant.
Calibration	The retrospective comparison of predicted outcomes to observed outcomes. Requires predicted outcomes. See "Decision Log," see "Adoption," see "Gap."
Connector	The integration layer between Sprig and your NPS / support platforms. Owns ingestion and tagging. Has been in "final testing" since Q4 2024.
Decision Log	The append-only record of every prioritization decision, its predicted outcome, and the actual outcome. Append-only in the technical sense. Empty in the practical sense.
Eligible PM	A PM whose team has adopted the framework and is past the 30-day onboarding period. There are currently 14 eligible PMs. 13 of them believe onboarding is still in progress.
Framework	The four-criterion process: Speed, Auditability, Adoption, Calibration. Also what everyone says they have before they describe a spreadsheet.

Glossary

P – T

TERM	DEFINITION
Predicted outcome	The author's stated expectation, recorded at decision time, used as the input to calibration. Hypothetically.
Prioritization rhythm	The team's regular cadence for revisiting and re-ordering work. Described as "weekly" in all documentation. Observed as "whenever there is an incident or a board meeting."
Retrospective	The 30-day review meeting where logged decisions are scored against observed outcomes. Sometimes held on day 45. Usually runs 90 minutes over schedule. Produces action items that are reviewed at the next retrospective.
Signal	Any qualitative or quantitative input to a decision. Formally: survey response, NPS comment, support ticket, sales call note. Informally: whatever the last person who spoke to a customer told the product team in a Slack message at 9 PM.
Sprig	The micro-survey product used by Tool D. The recommended solution. See "Conflict of Interest," which does not appear in this glossary.
Tool D	The recommended option. The tool built by the authors of this evaluation. The best tool by the criteria the authors defined. We are noting this again because the footnote on the table slide is small.

Accent replaces the rainbow stripe with a single editorial colour

It tints the heading and composes with `dark` — where `accent.dark` restores the stripe the spectrum suppression removed. The framework has a modifier for this. The framework team is very thorough.