

DECISION FRAMEWORK • Q3 2025

From Signal to Strategy

A 78-slide answer to the question "have you considered writing things down"

What this deck covers, in order

01 Why This Exists — slide 3

→ **The Framework — slide 7**

03 The Evaluation — slide 22

04 The Build — slide 31

05 The Results — slide 47

CONTEXT • COMPETITIVE DYNAMICS

The window for differentiation is narrowing

Commoditized infrastructure, compressed release cycles, and rising switching costs have cut the average durable advantage from 36 months to under 14. This slide will appear in every deck for the next two years regardless of whether that stays true.

“

The signal was always there. We just didn't have a system that forced us to look at it before we'd already decided.

”

— Head of Product, Pilot Team 3, in a retrospective where we then decided what we'd already decided

IMPACT · PILOT RESULTS

Six months of results across four product teams

SAME TEAMS, SAME CONDITIONS, SAME SPREADSHEET.



CALIBRATION RESULT • 6-MONTH PILOT

14X

Return on signal investment — on a
baseline the framework team defined.

Verified by nobody.

SECTION 01 • THE FRAMEWORK

We built a four–
component scoring
system. Two of the four
are used regularly


SIGNAL DEFINITION • WORKSHOP 04

Before we score signals, we need to agree on what a signal is

Three workshops in. The definition gets
socialized in a fifth before anyone can use it.

How signals move from input to decision

Four-stage pipeline — 11 weeks in, still in pilot



The framework has four components

Signal Intake

Everyone agrees it's a good idea. Nobody does it the week of the retrospective.

Scoring Model

Reconfigured by the head of product until the output agrees with the roadmap.

Decision Log

18 entries. Roughly 340 decisions.

Calibration Loop

The meeting exists. The predictions rarely do.

Signal Intake produces three outputs

1

Weekly Signal Brief

Top 10 signals, every Monday — unread, in a folder called "Framework Stuff."

2

Anomaly Alerts

2σ crossings on a 4-hour SLA. The PM replies in 4 hours, to ask what 2σ means.

3

Monthly Signal Index

Required reading before every retrospective. Nobody has read it. It is comprehensive.

The three things the framework connects

Signal

A verbatim, a metric move, a competitor's move. Or "what the VP heard at a conference."

Decision

A signal plus a deadline. Usually the signal is "we discussed it at the offsite."

Outcome

Compared to the prediction at retrospective. 18 logged; ~340 occurred.

Two failure modes the framework is designed to prevent

False signal amplification.

One loud voice — a customer, an analyst, a VP with a feeling. Capped at 30% of weight. Unless it's the CEO, in which case the cap is a guideline.

Signal hoarding.

Signals collected, decisions unlogged, so calibration learns nothing. The rule was a poster in the meeting room. It's now a free-pizza flyer.



Scoring Model Deep Dive

What the scoring model actually does

The most configurable component — a feature or a warning sign, depending on your team. Three dimensions:

- 1

Confidence
Independent corroborating sources, 1–5. Enterprise customers count as 1, regardless of volume.
- 2

Recency
Time-decay, configurable half-life. Set to two weeks, then surprised only recent news scores.
- 3

Strategic Relevance
Owner-scored, 1–5. A 5 correlates with whoever is presenting the roadmap this quarter.

The six signal dimensions, what they measure, and how they are scored

01	Confidence	Independent sources corroborating the signal	1-5 · AUTO · ENTERPRISE ALWAYS GETS A 4
02	Recency	Time-decay from signal date, configurable half-life	0.0-1.0 · AUTO · "SHORT" AFTER A BAD QUARTER
03	Relevance	Alignment to current strategic bets, owner-scored	1-5 · MANUAL · WHAT THE PM ALREADY PLANNED
04	Reach	Customers or segments affected	1-5 · AUTO · 5 IF THE ENTERPRISE ASKED
05	Effort	Engineering and design cost to act on the signal	1-5 · MANUAL · ENGINEERING'S EYEBROWS GO UP
06	Confidence delta	Change in confidence since the last scoring cycle	-5 TO +5 · AUTO · DID ANYONE TALK TO A CUSTOMER

Scoring model: before and after the calibration loop

BEFORE CALIBRATION

Equal weights, 33% each. Honest that we're guessing.



AFTER CALIBRATION

Weighted by your team's history — three months of it, from a quarter everyone calls atypical.

The shift takes two cycles — 60 days, or 14 months, depending on who you ask.

Two intake modes for different signal types

Structured Intake

NPS, tickets, win/loss notes. 94% of the data, 12% of the decisions.

Unstructured Intake

Hallway talk, conferences, a board member's dinner. 6% of the data, 88% of the decisions.

How a decision moves through the framework



What the framework does not do

Doesn't make decisions — it formalizes the one you'd have made anyway.

Doesn't replace discovery — it routes whatever surfaces before the roadmap locks.

Doesn't work without the Decision Log. 18 entries; ~340 decisions made.

Doesn't guarantee alignment — the same argument, just earlier in the quarter.

Doesn't scale below P2 — the daily 90% still lives in Slack threads from 2022.

Four things that must be true before you begin

- 01 A monthly prioritization cadence. "Whenever someone escalates loudly" doesn't count — it just gets louder.
- 02 Someone who owns signal collection. Which means defining a signal first. See slide 8.
- 03 Leadership logging rationale, not just outcomes. This slide exists because that part is hard.
- 04 90 minutes a week. Budget four.

SECTION 02 WEIGHS FOUR TOOLS AGAINST FOUR
CRITERIA — DEFINED BY THE TEAM THAT BUILT
ONE. DISCLOSED IN A FOOTNOTE.

SECTION 01 OF 05 COMPLETE

The framework is
specified — the real
question is build or
buy

SECTION 02 · THE EVALUATION

We evaluated four
tools. One of them was
built by the evaluation
team

Four requirements every decision system must meet

01 **Speed**

Decisions close inside their window. Six months to calibrate isn't fast. We won't dwell on that here.

02 **Auditability**

Every decision above the threshold carries a traceable rationale — for after the launch goes badly.

03 **Adoption**

Weekly use, or calibration never runs. We budgeted 90 minutes per PM. Actual: 11.

04 **Calibration**

It has to improve. A static model is a spreadsheet with a dashboard nobody checks.

We evaluated four intake tools against the criteria

Tool A · Chorus

- ✓Speed
- ✖Auditability
- ✓Adoption
- ✖Calibration

Great call recording, no logging, no calibration. The sales team already uses it.

Tool B · Productboard

- ✖Speed
- ✓Auditability
- ✓Adoption
- ✖Calibration

Solid intake, no calibration. Setup was "3–4 weeks." It took 11.

Tool C · Notion

- ✓Speed
- ✓Auditability
- ✖Adoption
- ✖Calibration

Build anything. Teams built seven, then debated the canon for two quarters.

Tool D · Sprig + Decision Log

- ✓Speed
- ✓Auditability
- ✓Adoption
- ✓Calibration

Built by this evaluation's authors. Meets all four — as they defined them.

The four tools side by side

CRITERION	CHORUS	PRODUCTBOARD	NOTION	SPRIG + LOG
Speed	✓	✗	✓	✓
Auditability	✗	✓	✓	✓
Adoption	✓	✓	✗	✓
Calibration	✗	✗	✗	✓
Setup time	1 day	3–4 weeks	40+ hours	Same day

♦ Criteria defined by the team building Sprig + Log. We're transparent about it. It's in the footnote.

How we sort the four tools against our two axes

COVERAGE • COST

High coverage / Low cost

Sprig + Log — best coverage, lowest TCO. Built by the evaluation team.

Productboard — narrower, here to show we looked.

High coverage / High cost

Notion — full coverage in theory. We tried it: seven versions, a quarterly canon debate.

Low coverage / Low cost

Chorus — cheap, three criteria short. The sales team likes it.

Low coverage / High cost

None — either the signal or a gap. We're treating it as a signal.

Applying the criteria to the tools — here is where the evidence points

The evidence favors Tool D



Sprig plus a light Decision Log meets all four criteria inside the 90-minute budget — and was recommended before the evaluation began.


The path is not self-executing

Needs a connector to NPS and support. In review for 11 weeks. Still "week one."

The Decision Log is the hardest part

Not technical — cultural. Log before you close. A "habit change" in every deck since Q3 2023.

The evaluation came down to one question the vendors could not answer



The left card is struck through; the DECISION connector is bold. The conclusion came first — the slide was built to hold it.

We are building, not buying

DECISION · 2026 Q1

BUILD

Owens the policy, the loop, the timeline — plus on-call and every future "why did it score that."

WHY NOT BUY

None expose calibration weights.
The weights are the product.
That was the finding.

WHY NOT DELAY

The window closes in 18 months.
In the deck since Q1 2025.

How we make calls when the spec is silent

- 01 Default to the choice that's cheaper to reverse — unless reversing it needs a meeting.

- 02 Name the actor, never the system. The system can't be held accountable. The actor can be reorganized.

- 03 Write the bet on the slide with the choice. Not always where it gets reviewed.

What the build covers

Three phases, four workstreams. We own the policy, the loop, the timeline — and whatever Phase 2 becomes.

1

Phase 01 — Architecture

Policy live. Log accepting entries. One pilot team.

2

Phase 02 — Operations

Multi-team calibration. Automated weight updates. Board-ready data.

3

Phase 03 — Scale

Org-wide enablement. In the roadmap since 2024.

P

SECTION 03 • THE BUILD

Phase 01 through
Phase 03 ships the
architecture, the
operations, and the
scale

What ships in each phase, by workstream

WORKSTREAM	PHASE 01	PHASE 02	PHASE 03
Signal Intake	Connector v1	Multi-source dedupe	Anomaly auto-routing
Scoring	Equal-weights model	Per-team calibration	Per-decision profiles
Decision Log	Append-only schema	Outcome auto-pairing	Examiner export
Adoption	One pilot team	—	Org-wide enablement

Phase 3 holds org-wide enablement. In the roadmap since 2024. Phase 2 is ongoing.

Who owns each part of the framework lifecycle

Signal custody

Signal owner

Manages intake quality. Tunes weights only by choosing what to surface.

Policy

Framework operator

Owns scoring, cadence, rollback. One person. Noted in the risk register.

Consumption

Product team

Runs intake and logging. Mostly asks the operator to change the weights.

Oversight

Auditor

Reads the audit trail. Read it once, found 18 entries, asked if that was expected.

Three phases get us from decision to org-wide adoption

PHASE 01

Architecture

Scopes build, buy, defer.
Output: a decision record that
needs its own architecture
phase.



PHASE 02

Pilot

One team, one decision type,
one quarter. Ends when the
retrospective happens once.



PHASE 03

Rollout

Five teams in two months.
Planned for Q2 — three years
running.

Three milestones mark Phase 01 complete

MILESTONE A

Scoring policy in production

Signed policy runs end-to-end.
Leadership's first reply: can the weights be adjusted? They can.

MILESTONE B

Per-team weights

Distinct weights per team, no forks. Each team still wants its own.

MILESTONE C

Per-decision-class profiles

A profile per class takes minutes. Agreeing what a class is takes a workshop. See slide 8.

Phase 01 acceptance — what shipped, what slipped, what stayed open

✓	Scoring policy live across all pilot teams	
✓	Decision Log audit trail readable by Auditor role	shipped 2026-Q1
✓	One reference team running weekly cadence	one team
—	Examiner pack auto-generation	was Phase 01 · now Phase 1b
○	Adoption above 90%	Phase 03
○	Culture change	not in roadmap

Decisions used to require a quarterly re-litigation

BEFORE AND AFTER THE FRAMEWORK

BEFORE

First principles every time. Close: 4 hours, p99 an offsite. Outcome: what the senior person wanted, called consensus.



AFTER

Resolved against logged weights. Close: 18 minutes. Outcome: what the model says — once the weights match the senior person.

The architecture change is the calibration loop. The culture change is still in Phase 02.

The same transition for the board deck

BEFORE AND AFTER · BANNER-TAG MODIFIER

BEFORE

First principles, 4 hours, no trail, no calibration. We called it "moving fast."

→

AFTER

Logged, weighted, calibrated. The trail has 18 entries. We call it "moving thoughtfully."

The same build decision in banner-tag chrome

DECISION · BANNER-TAG MODIFIER

BUILD

Owns everything the rejected path said "a vendor would own."

WHY NOT BUY

None expose calibration weights — the criterion set by the team that builds frameworks.

WHY NOT DELAY

18-month window. Month 7.
Phase 02 planned for month 6.
We're aware.

Recalibration used to require a coordinated freeze

BEFORE — MANUAL RECALIBRATION

Schedule a window, freeze decisions, swap weights, verify, unfreeze. 18 hours. Last year's way, fine until this deck.

AFTER — VERSION-FLOOR RECALIBRATION

The loop emits a versioned policy; teams pick it up on refresh. No freeze. Hasn't shipped yet.



How to roll this out across your organization

STEP 01

Pick one team and one decision type

Start with a team that has a real rhythm. No rhythm, no framework.

STEP 02

Log everything, decide nothing differently

Just log, for a month. "Low-effort." Highest dropout rate.

STEP 03

Run your first retrospective

Day 30, score against outcomes. No outcomes? Score the room, call it "qualitative."

STEP 04

Expand to a second team

Your evidence: one team ran one retrospective. Use it confidently.

The weekly practice in three moves

STEP 01

Sense

Observed inputs, never invented. Skipped about 70% of the time.

STEP 02

Score

A signal is data once it carries a number. Calibrating it takes 14 months.

STEP 03

Decide

A signal plus a deadline. The loop closes if anyone logged a prediction.

The case for the framework in three moves

Claim

Calibrated prioritization, audit-grade custody. No more re-litigating every quarter.

Evidence

Six months, four teams.
Calibration ran once. The four teams are now three. None in the baseline.

Implication

The framework works. The deck says so. The deck's authors are the framework team.

Wiring a signal into the framework is three lines of application code

JAVASCRIPT · DECISIONFRAMEWORK SDK V2 · 847 TRANSITIVE DEPENDENCIES

```
import { DecisionFramework } from "@company/signal-sdk";

const framework = new DecisionFramework({ configFile: "./framework.config.json" });

// Score a signal at intake
const score = await framework.score(signal, { dimensions: ["confidence", "recency", "relevance"] });

// Log every decision – calibration depends on it
// (nobody calls this line in production, but it is here)
const entry = await framework.decisions.log(decision, { signals: [signal.id], rationale });
```

BEFORE & AFTER • SCORING MECHANICS

Spreadsheet-driven scoring versus framework-driven scoring that is basically also a spreadsheet

BEFORE • THE HONEST SPREADSHEET

```
# Manual scoring. Auditable in the  
# sense that you can see who last  
# edited the file  
import pandas as pd  
  
signals = pd.read_csv('./signals.csv')  
signals['score'] = signals.apply(  
    lambda r: 0.33*r.confidence + 0.33*r.recency + 0.33*r.relevance,  
    axis=1,  
)  
signals.to_csv('./scored.csv')
```

AFTER • THE FRAMEWORK

```
# Calibrated weights, signed policy,  
# every score is audit-logged,  
# same math as the spreadsheet  
from decision_framework import Calibrator  
  
calibrator = Calibrator.load('./policy.json')  
for signal in calibrator.intake.unscored():  
    calibrator.score(signal)  
    calibrator.decisions.log_if_relevant(signal)
```

SECTION 04 PRESENTS SIX MONTHS OF
PRODUCTION DATA. THE DENOMINATOR IS
CAREFULLY CHOSEN.

SECTION 03 OF 05 COMPLETE

The framework is
built. Twelve percent
of eligible PMs use it

SECTION 04 • THE RESULTS

Six months of data.
Eighteen logged
decisions. One
calibration cycle

The pilot team measured what the pilot team built.

Before the results, what you should have learned in the first 46 slides

- | | |
|----|--|
| 01 | Four components. Two used regularly. → Section 01 |
| 02 | The evaluation team recommended the tool the evaluation team built. → Section 02 |
| 03 | Phase 01 shipped. Phase 03 in the roadmap since 2024. → Section 03 |
| 04 | Adoption is 12%. Target 90%. The gap is "cultural." → this section |
| 05 | The calibration loop has run once. We call it "calibrated." → this section |

Where we are against quarter targets



18 min

p99 decision close

TARGET 20 MIN, BEATING TARGET

18

Decisions logged

TARGET 340, GAP IS "CULTURAL"


1

Calibration cycles run

TARGET 6, GAP IS "STRUCTURAL"

Phase 01 readiness, by workstream


The status pills reflect the most optimistic reading of the data.



Source: Linear · "blocked" means blocked, we are working on the wording


The same data, dark canvas

Same bars, same numbers. Dark makes the 12% adoption bar feel intentional, not alarming.
That is not why we have a dark modifier. It is why this slide uses it.



Same data, minimal treatment

The lucent strip is gone, so the adoption number is louder. An argument for the non-minimal treatment.




Source: Linear · refreshed 2026-05-07

H1 2026 · 1,840 PERSON-HOURS

Where the engineering quarter went

The "Toil and on-call" wedge is the one nobody put in the roadmap.



Last updated 2026-05-07 · the 9% is probably higher

FRAMEWORK ARC

How the framework arrived in production

Four stages over eighteen months. The connective tissue is described as "momentum."



Cross-functional sign-off · "cross-functional" means two teams instead of one

FINDING 01 • STRUCTURED INTAKE

Structured intake performed above expectations — volume and latency were not the problems we thought

What worked

94% of structured signals, no manual touch, 4-minute latency. The part of the demo we show.

What required tuning

18% NPS-classification error for two weeks. In the appendix, not the headline.

KEY INSIGHT

Viable as designed. The headline numbers are accurate. The denominator is carefully chosen.

Three scoring failure modes found in the pilot

1

Recency dominance

Recency set above 50% — last week felt more real than six months of NPS. Capped at 40%, until a VP overrides it.

2

Source concentration

One customer was 34% of intake. Also 31% of revenue, so the diversity floor waited for renewal.

3

Outcome misclassification

"Improve retention" won't score. "Make customers happier" was submitted twelve times, scored 5 each.

WALK THESE QUESTIONS WITH ME IN 60-90
MINUTES. THE OUTPUT IS A DESIGN WE CAN
EXECUTE — OR AGREEMENT THAT WE NEED
ANOTHER SESSION TO DESIGN IT.

WHAT WOULD HELP US MOVE FORWARD

Next step is a
working session, not
a debate

SECTION 05 • THE PATTERN LIBRARY

The story closed on slide 57. These slides demonstrate authoring patterns

The reference appendix: the layout variants and modifiers used above. Its counter runs independent of the dark divider's — so it stamps 01, not 05.

DARK VARIANT • CONTENT

The token system handles dark without per-element overrides

Every colour is a custom property that remaps on *dark* — cards, text, borders, the suppressed spectrum bar. Popular with executives who want "more premium." Same content, darker box.

HEADER AND FOOTER

Header stays uppercase — footer renders as written

Set `header:` and `footer:` in frontmatter. The header always uppercases; the footer renders as written — which is how "FINAL v3 USE THIS ONE" reaches the board.

The card stack renders cleanly on dark backgrounds

Cards fill `--bg-alt`, borders `--border` — both remap on dark. Most decks don't. Theirs with break.

The accent border `--accent`, unchanged, and reads on dark. Chosen partly for that. uses You're welcome.

Body text `--text-body` — warm light, not pure white. Pure white on dark is readable and is unkind.

Two-card layouts work equally well inverted to dark

**One calibration question: what proves the model is improving, and what does a quarter without that proof cost?
In practice, "we'll revisit in Q3."**


Claim, then support — unchanged on dark. What changes is the screenshot your VP forwards with "see, this is what I mean."

LAYOUT · IMAGE

Image right is the default — text leads, evidence follows

For slides where the visual supports the argument.
Native aspect ratio, no cropping — usually a stock lighthouse called "a metaphor for signal."






MODIFIER · IMAGE MIRROR

Mirror flips the image to the other half — alias of legacy `image left`

`mirror` is the canonical orientation flag. The `left` alias survives for the 40 decks nobody wants to find.



Signal Pipeline · Reference Visualization

Weekly Signal Brief — shipped every Monday, opened by an estimated 3 of 14 eligible PMs

Code chips render inside card titles and body alike

Signal Intake

v2.4

94% of structured signals. The other 6% are CEO Slack messages, escalated unscored.

Scoring Model

configurable

Defaults `33 / 33 / 33`. Most set recency to `99` after a bad quarter, call it calibrated.

Decision Log

required

Every change above `P2` needs a rationale. Introduced Q2 2024. Not yet enforced.

Calibration Loop

monthly

First real update after `2 cycles` — about `14 months`.

A trailing italic-only paragraph becomes an annotation

Signal Intake

Weekly collection across conversations and market data.

Scoring Model

Three dimensions, configurable by seven admins.

Decision Log

Every decision recorded. Enforced by honor system.

Calibration Loop

Predictions vs outcomes. Averages 2.3 attendees.

KEY INSIGHT

The calibration loop separates teams that learn from teams that repeat the same mistakes more thoroughly.

✦ *Pilot retrospective: six months, four teams, one since reorganized out of existence.*

A trailing blockquote becomes a key insight; a plain paragraph becomes a below-note

Signal Intake

The one everyone demos.

Scoring Model

The one everyone reconfigures.

Decision Log

The one nobody fills out.

Calibration Loop

The one that finds no predictions.

KEY INSIGHT

The calibration loop separates teams that learn from teams with very thorough meeting notes about not learning.

Trailing blockquote → key insight; trailing paragraph → below-note.

Compact tightens the spacing scale ~25%

What changes

`--sp-xs` through `--sp-2xl` shrink. Card gaps, gutters, and padding follow, via `var()`.

What does not change

Type, palette, and chrome reservation are untouched. A density flag, not a new layout.

When to reach for it

One card too many, prose a line long, or a denser rhythm.

Composition

Composes with `dark`, `accent`. Not with `title`, `divider`, `image full`.

MODIFIER • LOOSE

Loose is the inverse — more breathing room, same layout machinery

The spacing scale grows ~25%. Reach for *loose* when a slide carries one editorial point and wants to feel quiet.

KEY INSIGHT

Density is not importance. **loose** says this page deserves room — because it carries one thing well. The framework team applied it to this quote, which the framework team wrote about the framework.

Routing signals by source type versus routing by confidence tier

ROUTE BY SOURCE TYPE

Structured to Intake, unstructured to the owner.
Clean — until an enterprise customer hits an edge.



ROUTE BY CONFIDENCE TIER

Above 3.0 to scoring, below to review. Adaptive, once confidence is calibrated. We're in month three, calling it cycle one.

Both paths land on the same PM, who asks why the framework keeps routing things to them.

Chosen flags the right-hand card as the winner

QUARTERLY RECALIBRATION

Reviewed once a quarter. Simple, blind to everything between. What we did before — undiscussed here.



CONTINUOUS CALIBRATION

Updates every retrospective past 12 logged decisions. We have 18. The loop ran once.

The right card carries the accent edge and tint — the chosen path. Chosen after the choice was made.

Mirror composes with chosen — the accented card reads from the left

RETROSPECTIVE-TRIGGERED RECALIBRATION

Updates past 12 decisions. `mirror` left, `chosen` accents. Conclusion first — useful when they won't ask the sample size. It's 18.



QUARTERLY RECALIBRATION

Once a quarter, no outcome data. Source order preserved — the first card is always the alternative. Mirror flips rendering, not intent.

The below-note holds the caveat. The loop ran once. We call it continuous.

Mirror puts the hero card on the right; sub-cards stack on the left

The scoring model requires the loop

No calibration, static weights — a spreadsheet. Sub-cards go left, written to make the hero inevitable.

The Decision Log enables the loop

Loop needs outcomes; outcomes need logged decisions.
18 entries, one run. The card the presenter skips.

The calibration loop is what closes the framework

The only component that makes the model improve instead of documenting failures better. `mirror` puts it right — premise before conclusion.

Four switches to 4 columns; pair with compact for visual balance

Sense

Observe. Write it down.
Don't interpret yet.

Score

A signal is data once it
has a number.
Calibrated. Eventually.

Decide

A signal plus a deadline. A
recurring meeting is not a
deadline.

Review

The retrospective. The
loop closes here — if it
had data.

Glossary

A – F

TERM	DEFINITION
Adoption	PMs logging within 24h of a close. 5.3% now; target 90%. The gap is "cultural."
Auditability	Reconstruct any decision months later without its author. The author is always there anyway.
Calibration	Predicted vs observed. Requires predictions. See "Decision Log," "Adoption," "Gap."
Connector	Sprig ↔ your NPS / support platforms. In "final testing" since Q4 2024.
Decision Log	The append-only record of every decision. Append-only in theory, empty in practice.
Eligible PM	Past 30-day onboarding on an adopted team. 14 exist; 13 think they're still onboarding.
Framework	Speed, Auditability, Adoption, Calibration. Also what people call a spreadsheet.

Glossary

P – T

TERM	DEFINITION
Predicted outcome	The expectation logged at decision time, fed to calibration. Hypothetically.
Prioritization rhythm	Documented "weekly." Observed "whenever there's an incident or a board meeting."
Retrospective	Day-30 scoring of decisions vs outcomes. Often day 45. Always runs long.
Signal	A survey, an NPS comment, a ticket. Or a 9 PM Slack from whoever last saw a customer.
Sprig	The micro-survey behind Tool D. See "Conflict of Interest" — not in this glossary.
Tool D	The recommended option, built by this evaluation's authors, best by their criteria.

MODIFIER • ACCENT

Accent swaps the rainbow stripe for one editorial colour

It tints the heading and composes with `dark`, where `accent.dark` restores the suppressed stripe. The framework team is very thorough.