DECISION FRAMEWORK · Q3 2025

From Signal to Strategy

A 78-slide answer to the question "have you considered writing things down"

What this deck covers, in order

CONTEXT · COMPETITIVE DYNAMICS

The window for differentiation is narrowing

Three forces — commoditized infrastructure, compressed release cycles, rising switching costs — have cut the average durable advantage from 36 months to under 14. Teams that can't tell signal from noise in that window miss the timing. This slide will appear in every deck for the next two years regardless of whether the window keeps narrowing.

> “*The signal was always there. We just didn't have a system that forced us to look at it before we'd already decided.*”

— Head of Product, Pilot Team 3, in a retrospective where we then decided what we had already decided

IMPACT · PILOT RESULTS

Six months of results across four product teams

MEASURED AGAINST PRE-FRAMEWORK BASELINE, SAME TEAMS, SAME CONDITIONS, SAME SPREADSHEET.

73%
FASTER DECISION CLOSE

4.2×
SIGNAL RECALL

18
DECISIONS LOGGED

91%
TEAM ALIGNMENT

CALIBRATION RESULT · 6-MONTH PILOT

14X

Return on signal investment — calculated by the team that built the framework, using a baseline they defined. Independently verified by nobody.

SECTION 01 · THE FRAMEWORK

We built a four-component scoring system. Two of the four are used regularly

SIGNAL DEFINITION · WORKSHOP 04

Before we score signals, we need to agree on what a signal is

Three workshops so far. We are now in the fourth. Its output is a shared definition, to be socialized in a fifth workshop before anyone can use it.

ARCHITECTURE · SIGNAL PIPELINE

How signals move from input to decision

Four-stage processing pipeline — 11-week implementation, still in pilot

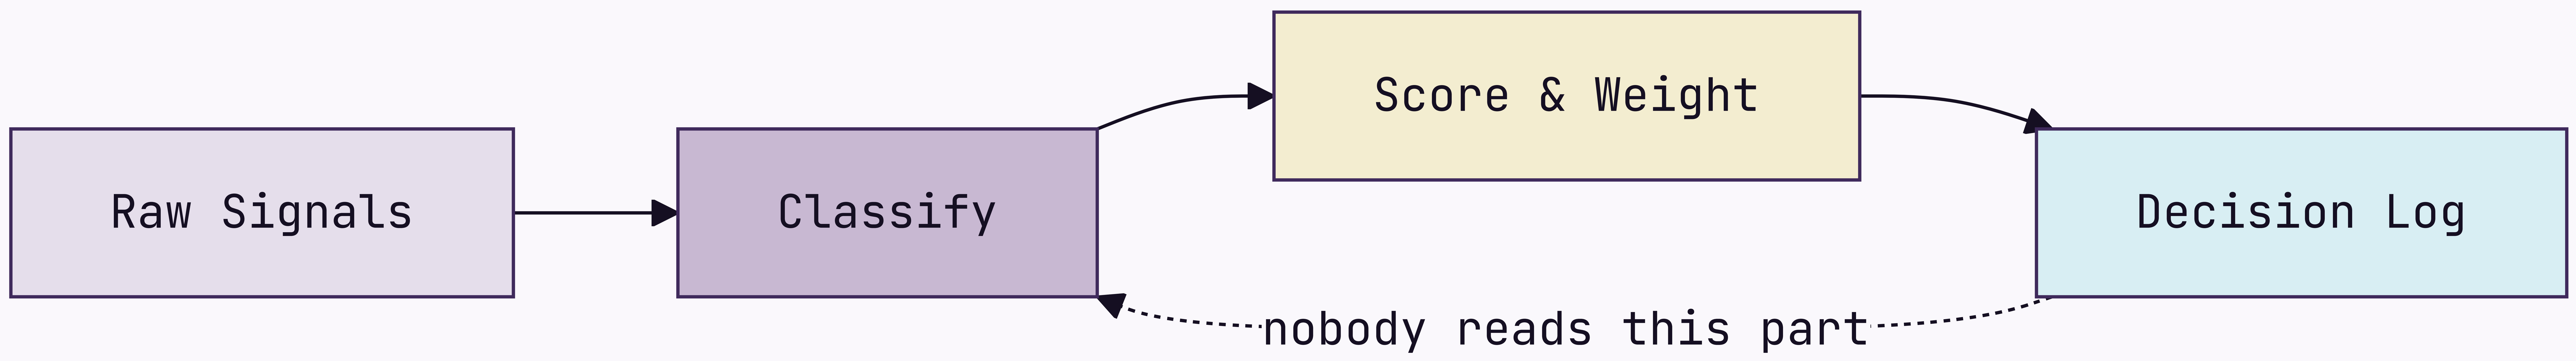

The framework has four components

Signal Intake

Weekly structured collection across customer conversations, market data, and competitive moves. Everyone agrees it's a good idea; nobody does it the week of the retrospective.

Scoring Model

Each signal scored on confidence, recency, and strategic relevance. Weights are team-configurable, which means the head of product reconfigures them until the output agrees with their roadmap.

Decision Log

Every decision recorded with its signals and criteria. Required artifact. 18 entries so far, against roughly 340 decisions.

Calibration Loop

Monthly retrospective comparing predicted outcomes to actual ones. The meeting exists. The predictions rarely do.

Signal Intake produces three outputs

1

Weekly Signal Brief

Last week's top 10 signals, sent to product leads every Monday — where it sits unread in a folder called "Framework Stuff."

2

Anomaly Alerts

Real-time flags when a signal crosses 2σ, routed to the accountable PM on a 4-hour SLA. The PM responds within 4 hours, to ask what 2σ means.

3

Monthly Signal Index

The source of truth for the calibration loop, required before each retrospective. Nobody has read it. It is comprehensive.

The three things the framework connects

Signal

The observed input — a verbatim, a metric move, a competitor announcement. Frequently confused with "things the VP heard at a conference."

Decision

A signal plus a deadline, logged with its rationale. In practice the signal is often "we discussed it at the offsite."

Outcome

The observed result, compared to the prediction at retrospective. 18 logged so far. Roughly 340 have occurred.

Two failure modes the framework is designed to prevent

False signal amplification.

A single loud voice — one enterprise customer, one analyst report, one VP with a feeling — dominates the decision. The model caps any single source at 30% of total weight. Unless that source is the CEO, in which case the cap is a guideline.

Signal hoarding.

Teams collect signals but log no decisions, so calibration has nothing to learn from. The rule — no log, no change above P2 — was printed on a poster in the meeting room. The poster has been replaced by a free-pizza flyer.

S

SECTION 01 · CONTINUED

Scoring Model Deep Dive

Dark panel + content · split-panel watermark

What the scoring model actually does

The scoring model is the most configurable component — a feature or a warning sign, depending on whether you trust your team not to set every dimension to 100. Three dimensions:

1. **Confidence**

   How many independent sources corroborate the signal. Ranges 1–5. Enterprise customers count as 1, regardless of volume.

2. **Recency**

   Time-decay from signal date; half-life is configurable. Most teams set it to two weeks, then wonder why they only act on recent news.

3. **Strategic Relevance**

   Owner-scored, 1–5, written justification required above 4. Scores of 5 correlate remarkably with whoever is presenting the roadmap this quarter.

The six signal dimensions, what they measure, and how they are scored

01	**Confidence**	Independent sources corroborating the signal	1-5 · AUTO · ENTERPRISE ALWAYS GETS A 4
02	**Recency**	Time-decay from signal date, configurable half-life	0.0-1.0 · AUTO · "SHORT" AFTER A BAD QUARTER
03	**Relevance**	Alignment to current strategic bets, owner-scored	1-5 · MANUAL · WHAT THE PM ALREADY PLANNED
04	**Reach**	Customers or segments affected	1-5 · AUTO · 5 IF THE ENTERPRISE ASKED
05	**Effort**	Engineering and design cost to act on the signal	1-5 · MANUAL · ENGINEERING'S EYEBROWS GO UP
06	**Confidence delta**	Change in confidence since the last scoring cycle	-5 TO +5 · AUTO · DID ANYONE TALK TO A CUSTOMER

Scoring model: before and after the calibration loop

BEFORE CALIBRATION

Equal weights — confidence, recency, and relevance each contribute 33%. Simple, and honest that we are basically guessing.

❯

AFTER CALIBRATION

Weights reflect your team's historical accuracy — except the team keeps changing and the history is three months from a quarter everyone agrees was atypical.

The shift from equal to calibrated weights takes two retrospective cycles — 60 days from adoption, or 14 months, depending on who you ask.

Two intake modes for different signal types

Structured Intake

Clear-schema signals: NPS verbatims, support tickets, win/loss notes. Ingested and scored automatically, zero manual handling. Produces 94% of the data and 12% of the roadmap decisions.

Unstructured Intake

No-schema signals: field observations, conference conversations, a board member who attended a dinner and has thoughts. Routed to the signal owner for manual classification. Produces 6% of the data and 88% of the roadmap decisions.

How a decision moves through the framework

1. **Signal Logged**

   Owner classifies and submits to the intake queue, if they remember

2. **Scored**

   Model applies the current weights — last updated in February

3. **Brief Published**

   Signal appears in the weekly brief that goes to product leads' spam

4. **Decision Logged**

   PM records rationale, options, predicted outcome. Optional in practice.

5. **Retrospective**

   Outcome scored against prediction. Predictions rarely exist. We improvise.

What the framework does not do

- It does not make decisions — it structures the information humans use to make the decision they were going to make anyway.
- It does not replace customer discovery — it scores and routes what discovery surfaces, if discovery surfaces anything before the roadmap locks.
- It does not work without the Decision Log — calibration needs outcome data. The log has 18 entries. We've made roughly 340 decisions.
- It does not guarantee alignment — it surfaces disagreement earlier, so you can have the same argument at the start of the quarter instead of the end.
- It does not scale down to individual features — below P2, the daily 90% of decisions is still made in Slack threads from 2022.

Four things that must be true before you begin

01 You have a regular prioritization cadence, monthly at minimum. If your cadence is "whenever someone escalates loudly," the framework just makes the escalation louder.

02 Someone owns signal collection as a primary responsibility — "primary" meaning more than 10% of their time, written into their job description. Which means agreeing what a signal is first. See slide 8

03 Leadership has agreed to log decisions with rationale, not just outcomes. This is harder than it sounds. This slide exists because it is harder than it sounds.

04 You have 90 minutes per week. Budget four.

SECTION 02 WEIGHS FOUR TOOLS AGAINST FOUR CRITERIA — DEFINED BY THE TEAM THAT BUILT ONE OF THEM. DISCLOSED IN A FOOTNOTE.

SECTION 01 OF 05 COMPLETE

The framework is specified — the real question is build or buy

SECTION 02 · THE EVALUATION

We evaluated four tools. One of them was built by the evaluation team

Four requirements every decision system must meet

01 **Speed**

Decisions must close inside the window they are relevant to. A framework that takes six months to calibrate is not fast. We will not address this directly on this slide.

02 **Auditability**

Every decision above the threshold carries a traceable rationale — for alignment, for compliance, and for reconstructing what happened after a launch goes badly.

03 **Adoption**

If the team won't use it weekly, calibration never runs. We put 90 minutes per PM in the deck. Actual usage is closer to 11.

04 **Calibration**

The system must improve over time. A static scoring model is a spreadsheet with extra steps and a dashboard nobody checks.

We evaluated four intake tools against the criteria

Tool A · Chorus

✓ Speed · – Auditability · ✓ Adoption · ✗ Calibration

Strong call recording, no decision logging, no calibration. Recommended by the sales team because they already use it.

Tool B · Productboard

✗ Speed · ✓ Auditability · ✓ Adoption · ✗ Calibration

Solid intake and prioritization, no calibration mechanism. Setup was "fine, 3–4 weeks." It took 11.

Tool C · Notion

✓ Speed · ✓ Auditability · – Adoption · ✗ Calibration

Flexible enough to build anything. Teams built seven versions and debated which was canonical for two quarters.

Tool D · Sprig + Decision Log

✓ Speed · ✓ Auditability · ✓ Adoption · ✓ Calibration

Built by the team that wrote this evaluation. Meets all four criteria as the team defined them. Recommended.

The four tools side by side

CRITERION	CHORUS	PRODUCTBOARD	NOTION	SPRIG + LOG
Speed	✓	✗	✓	✓
Auditability	✗	✓	✓	✓
Adoption	✓	✓	✗	✓
Calibration	✗	✗	✗	✓
Setup time	1 day	3–4 weeks	40+ hours	Same day

✦ *Evaluated against criteria defined by the team building Sprig + Log. We want to be transparent about this. It is in the footnote.*

How we sort the four tools against our two axes

COVERAGE · COST

High coverage / Low cost

Sprig + Log — best coverage, lowest TCO. Built by the evaluation team.
Productboard — narrower coverage, also here to show we looked.

High coverage / High cost

Notion build-out — full coverage in theory, premium upkeep. We tried it: seven versions, a quarterly debate over the canon.

Low coverage / Low cost

Chorus — cheap, three criteria uncovered, popular because the sales team already uses it.

Low coverage / High cost

None — and that is the signal. Or a gap in our vendor universe. We are treating it as a signal.

Applying the criteria to the tools — here is where the evidence points

The evidence favors Tool D

Sprig plus a lightweight Decision Log meets all four criteria inside the 90-minute weekly budget, reaches production the week it is adopted, and was recommended before the evaluation began. The evaluation confirmed this.

The path is not self-executing

Sprig needs a connector to your NPS and support platforms. It has been in review for 11 weeks. We still call it "week one."

The Decision Log is the hardest part

Not technically — culturally. PMs must log decisions before they close. Called a habit change in every deck since Q3 2023. It remains one.

The evaluation came down to one question the vendors could not answer

BUY A VENDOR FRAMEWORK

Three vendors evaluated. None expose calibration weights to the customer — the criterion that eliminated all three. It was added after the evaluation team decided to build.

❯

BUILD THE FRAMEWORK IN-HOUSE

Owns the scoring policy, the calibration loop, the timeline. The window closes in 18 months; a vendor cutover would eat nine — which is how we knew building was faster before the evaluation began.

The left card is struck through. The DECISION connector is bold. The conclusion was reached before the slide was built. The slide was built so the conclusion would have a slide.

We are building, not buying

DECISION · 2026 Q1

BUILD

Owns the scoring policy, the calibration loop, the timeline — and the maintenance, the on-call rotation, and any future explanation of why the framework scored the wrong thing.

WHY NOT BUY

Three vendors, none expose calibration weights. The weights are the product. We can't buy the product without them. That was the finding.

WHY NOT DELAY

The window closes in 18 months. This sentence has been in the deck since Q1 2025.

How we make calls when the spec is silent

01 We default to the choice that is cheaper to reverse — unless reversing it would need a meeting, in which case we pick something and document it as a decision.

02 We name the actor, never the system. The system cannot be held accountable. The actor can be reorganized.

03 We write the bet on the same slide as the choice. We do not always write down where the bet gets reviewed.

What the build covers

Three phases, four workstreams. We own the policy, the loop, and the timeline — and whatever Phase 2 turns into.

1 **Phase 01 — Architecture**

Scoring policy live. Decision Log accepting entries. One pilot team in weekly cadence.

2 **Phase 02 — Operations**

Multi-team calibration. Automated weight updates. Before-after data you can defend in a board update.

3 **Phase 03 — Scale**

Org-wide enablement. In the roadmap since 2024. We remain committed.

P

SECTION 03 · THE BUILD

Phase 01 through Phase 03 ships the architecture, the operations, and the scale

What ships in each phase, by workstream

WORKSTREAM	PHASE 01	PHASE 02	PHASE 03
Signal Intake	Connector v1	Multi-source dedupe	Anomaly auto-routing
Scoring	Equal-weights model	Per-team calibration	Per-decision profiles
Decision Log	Append-only schema	Outcome auto-pairing	Examiner export
Adoption	One pilot team	—	Org-wide enablement

Phase 3 is where org-wide enablement lives. It has been in the roadmap since 2024. Phase 2 is ongoing.

Who owns each part of the framework lifecycle

Signal custody
Signal owner

Manages intake quality and source diversity. Never tunes weights directly — only indirectly, by choosing which signals to surface.

Policy
Framework operator

Owns scoring policy, calibration cadence, version floors, and rollback. Currently one person. Noted in the risk register.

Consumption
Product team

Runs intake and decision-logging. Mostly interacts with the framework by asking the operator to adjust the weights.

Oversight
Auditor

Reads the Decision Log audit trail; cannot edit weights. Has read it once. Found 18 entries. Asked if that was expected.

Three phases get us from decision to org-wide adoption

PHASE 01

Architecture

Scopes what we build, buy, and defer. Output is an architecture decision record — which will itself need an architecture phase before it can be approved.

PHASE 02

Pilot

One team, one decision type, one quarter. Ends at production cadence, meaning the retrospective happened at least once.

PHASE 03

Rollout

Five teams in two months, ending above 90% adoption. Planned for Q2 — for three consecutive years.

Three milestones mark Phase 01 complete

MILESTONE A

Scoring policy in production

The signed policy runs end-to-end and the first calibrated brief lands in leadership's inbox. They reply asking whether the weights can be adjusted. They can.

MILESTONE B

Per-team weights

One framework carries distinct weights per team without forks; recalibration is a single update. Each team will still want their own.

MILESTONE C

Per-decision-class profiles

Authoring a profile for one decision class takes minutes. Agreeing what counts as a decision class takes a workshop series. See slide 8.

Phase 01 acceptance — what shipped, what slipped, what stayed open

- ✓ Scoring policy live across all pilot teams
- ✓ Decision Log audit trail readable by Auditor role — shipped 2026-Q1
- ✓ One reference team running weekly cadence — one team
- − Examiner pack auto-generation — was Phase 01 · now Phase 1b
- ○ Adoption above 90% — Phase 03
- ○ Culture change — not in roadmap

Decisions used to require a quarterly re-litigation

BEFORE AND AFTER THE FRAMEWORK

BEFORE

Every prioritization debate started from first principles. Average close: 4 hours. p99: an entire offsite. Outcome: whatever the most senior person wanted, expressed as consensus.

→

AFTER

Decisions resolve against logged weights and prior calibration. Average close: 18 minutes. Outcome: whatever the model outputs, after the weights are adjusted to reflect what the most senior person wants.

The architecture change is the calibration loop. The culture change is still in Phase 02.

The same transition for the board deck

BEFORE AND AFTER · BANNER-TAG MODIFIER

BEFORE

Every debate from first principles. Average close: 4 hours. No audit trail, no calibration. Decisions were made, outcomes happened, nobody connected them. This was called "moving fast."

→

AFTER

Decisions log their rationale. Scoring is weighted and calibrated. An audit trail exists; it has 18 entries. We are moving slower. We are calling this "moving thoughtfully."

The same build decision in banner-tag chrome

DECISION · BANNER-TAG MODIFIER

BUILD	WHY NOT BUY	WHY NOT DELAY
Owns the scoring policy, the calibration loop, and all the things the evaluation said "a vendor would own" in the rejected path.	Three vendors, none expose calibration weights. The evaluation team made "exposes calibration weights" a required criterion. The evaluation team builds frameworks.	The window closes in 18 months. We are in month 7. Phase 02 is planned for month 6. We are aware of this.

Recalibration used to require a coordinated freeze

BEFORE — MANUAL RECALIBRATION

Operators schedule a window, freeze new decisions, swap weights, verify, lift the freeze. Average pause: 18 working hours. Also known as the way we did it last year, which everyone agreed was fine until this deck was written.

❯

AFTER — VERSION-FLOOR RECALIBRATION

The loop emits a new policy with an incremented version; teams pick it up on next refresh. No freeze, no cutover. This is the good path. It is also the one that has not shipped yet.

How to roll this out across your organization

STEP 01

Pick one team and one decision type

Start with a team that already has a prioritization rhythm. If no team has one, the framework cannot help you. This is step one.

STEP 02

Log everything, decide nothing differently

In the first month, don't change how you decide. Just log. Described as low-effort. It has the highest dropout rate.

STEP 03

Run your first retrospective

At day 30, score the logged decisions against outcomes. If no outcomes were logged, score the energy in the room and call it "qualitative calibration."

STEP 04

Expand to a second team

With one retrospective done, you have evidence. The evidence is that one team ran one retrospective. Use it confidently.

The weekly practice in three moves

STEP 01

Sense

Observed inputs, never invented — and skipped about 70% of the time.

STEP 02

Score

A signal is data once it carries a number. Calibrating that number takes 14 months.

STEP 03

Decide

A decision is a signal plus a deadline. The loop closes if anyone logged a prediction.

The case for the framework in three moves

Claim

Calibrated prioritization with audit-grade decision custody. We stop paying re-litigation cost on every quarterly review.

Evidence

The pilot ran six months across four teams. Close-time dropped. Calibration ran once. The four teams have since been reorganized into three. None count in the "before" baseline.

Implication

The framework works. The deck says so. The deck was written by the framework team. The framework told us to trust the framework team.

IMPLEMENTATION · DECISION PIPELINE

Wiring a signal into the framework is three lines of application code

JAVASCRIPT · DECISIONFRAMEWORK SDK V2 · 847 TRANSITIVE DEPENDENCIES

```javascript
import { DecisionFramework } from "@company/signal-sdk";

const framework = new DecisionFramework({ configFile: "./framework.config.json" });

// Score a signal at intake
const score = await framework.score(signal, { dimensions: ["confidence", "recency", "relevance"] });

// Log every decision — calibration depends on it
// (nobody calls this line in production, but it is here)
const entry = await framework.decisions.log(decision, { signals: [signal.id], rationale });
```

BEFORE & AFTER · SCORING MECHANICS

Spreadsheet-driven scoring versus framework-driven scoring that is basically also a spreadsheet

BEFORE · THE HONEST SPREADSHEET

```
# Manual scoring. Auditable in the
# sense that you can see who last
# edited the file
import pandas as pd

signals = pd.read_csv('./signals.csv')
signals['score'] = signals.apply(
    lambda r: 0.33*r.confidence + 0.33*r.recency + 0.33*r.relevance,
    axis=1,
)
signals.to_csv('./scored.csv')
```

AFTER · THE FRAMEWORK

```
# Calibrated weights, signed policy,
# every score is audit-logged,
# same math as the spreadsheet
from decision_framework import Calibrator

calibrator = Calibrator.load('./policy.json')
for signal in calibrator.intake.unscored():
    calibrator.score(signal)
    calibrator.decisions.log_if_relevant(signal)
```

SECTION 04 PRESENTS SIX MONTHS OF PRODUCTION DATA. THE CALIBRATION LOOP HAS RUN ONCE. THE ADOPTION GAP IS DESCRIBED AS CULTURAL. THE DENOMINATOR IS CAREFULLY CHOSEN.

SECTION 03 OF 05 COMPLETE

The framework is built. Twelve percent of eligible PMs use it

SECTION 04 · THE RESULTS

Six months of data. Eighteen logged decisions. One calibration cycle

The pilot team measured what the pilot team built.

RECAP · SECTIONS 01 THROUGH 03

Before the results, what you should have learned in the first 46 slides

01	The framework has four components. Two are used regularly. → Section 01
02	The evaluation team recommended the tool the evaluation team built. → Section 02
03	Phase 01 shipped. Phase 02 is planned for Q2. Phase 03 has been in the roadmap since 2024. → Section 03
04	Adoption is at 12%. The target is 90%. The gap is described as cultural. → this section
05	The calibration loop has run once. We are calling this "calibrated." → this section

Where we are against quarter targets

94%

Signal-classification success

TARGET 99%, GAP IS "KNOWN ISSUE"

18 min

p99 decision close

TARGET 20 MIN, BEATING TARGET

18

Decisions logged

TARGET 340, GAP IS "CULTURAL"

1

Calibration cycles run

TARGET 6, GAP IS "STRUCTURAL"

H1 2026 · PHASE 01 READINESS

Phase 01 readiness, by workstream

Snapshot at 14:00 UTC. Status pills reflect the most optimistic reading of the data.

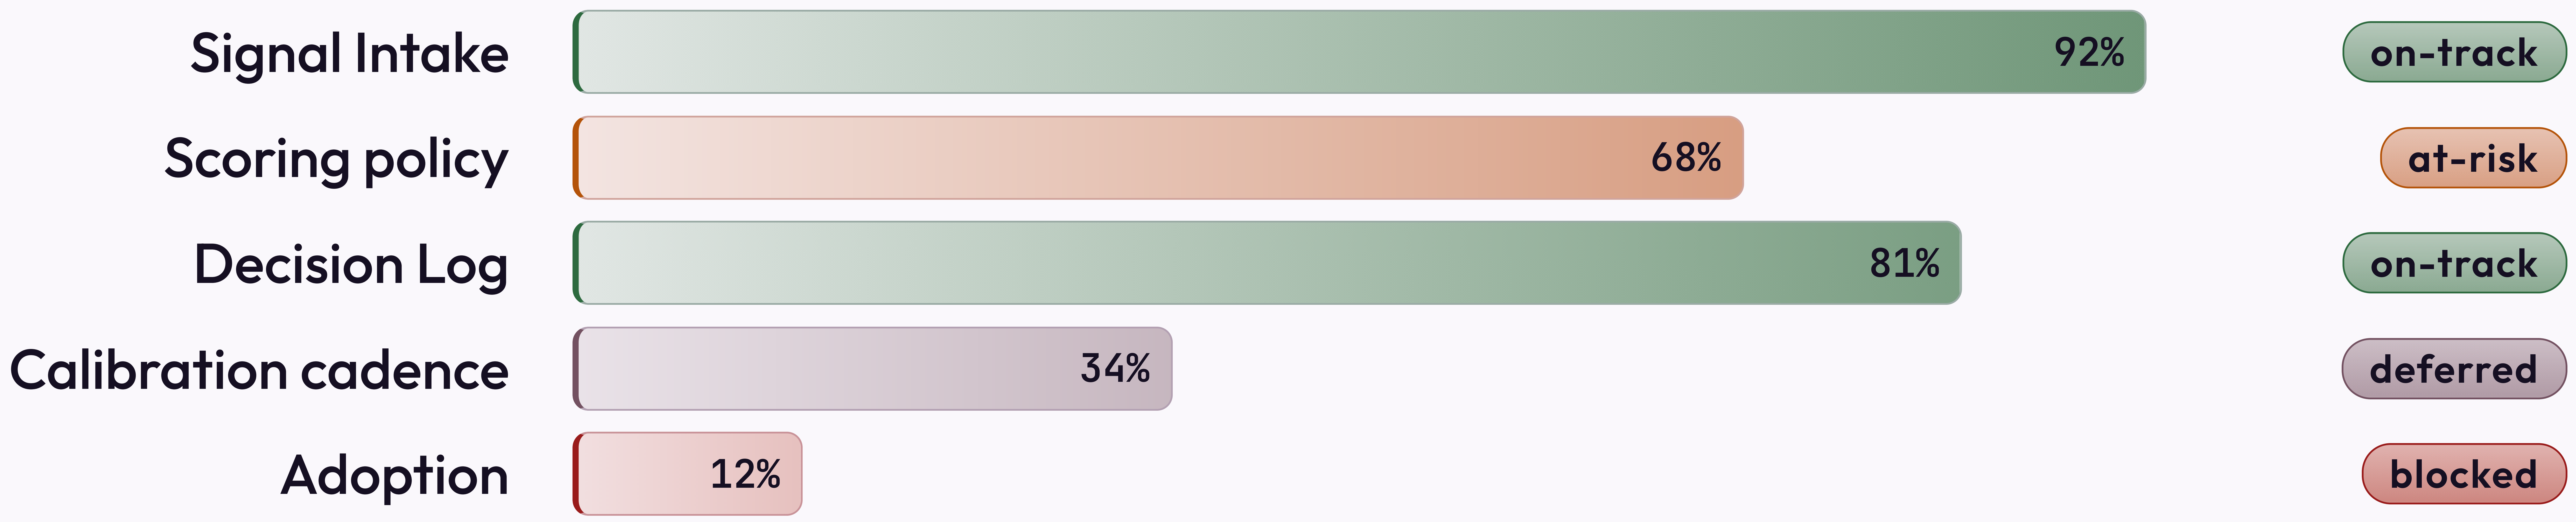

Source: Linear · refreshed 2026-05-07 · "blocked" means blocked, we are working on the wording

H1 2026 · PHASE 01 READINESS

The same data, dark canvas

The bars, pills, and numbers are identical. The dark canvas makes the 12% adoption bar feel more intentional and less alarming. This is not why we have a dark modifier. It is why this slide uses it.

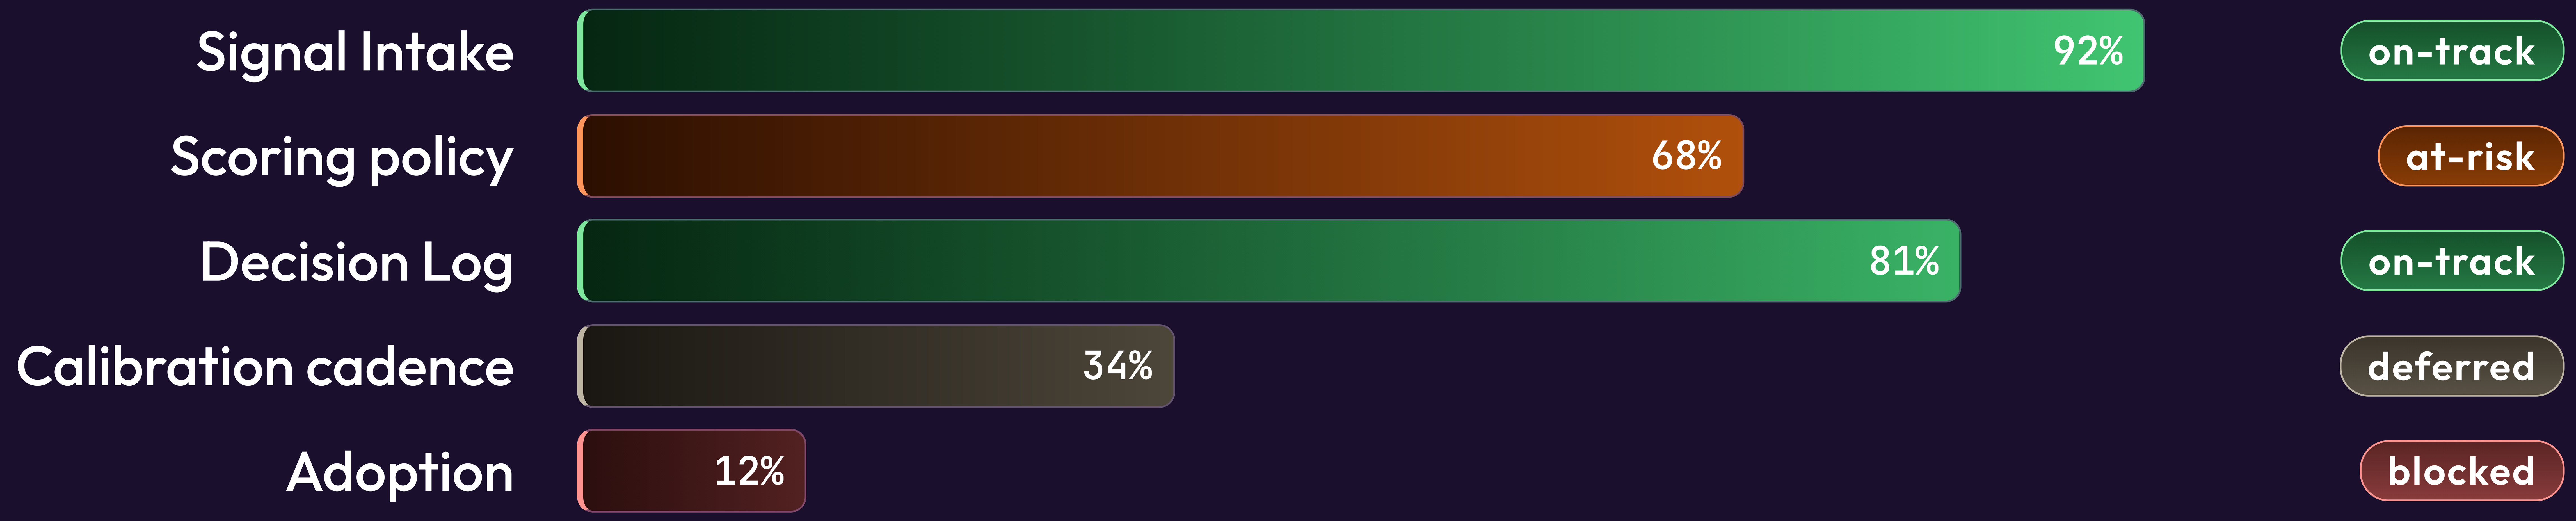

H1 2026 · PHASE 01 READINESS

Same data, minimal treatment

The lucent strip is gone. The adoption number is more visible. This is an argument for the non-minimal treatment.

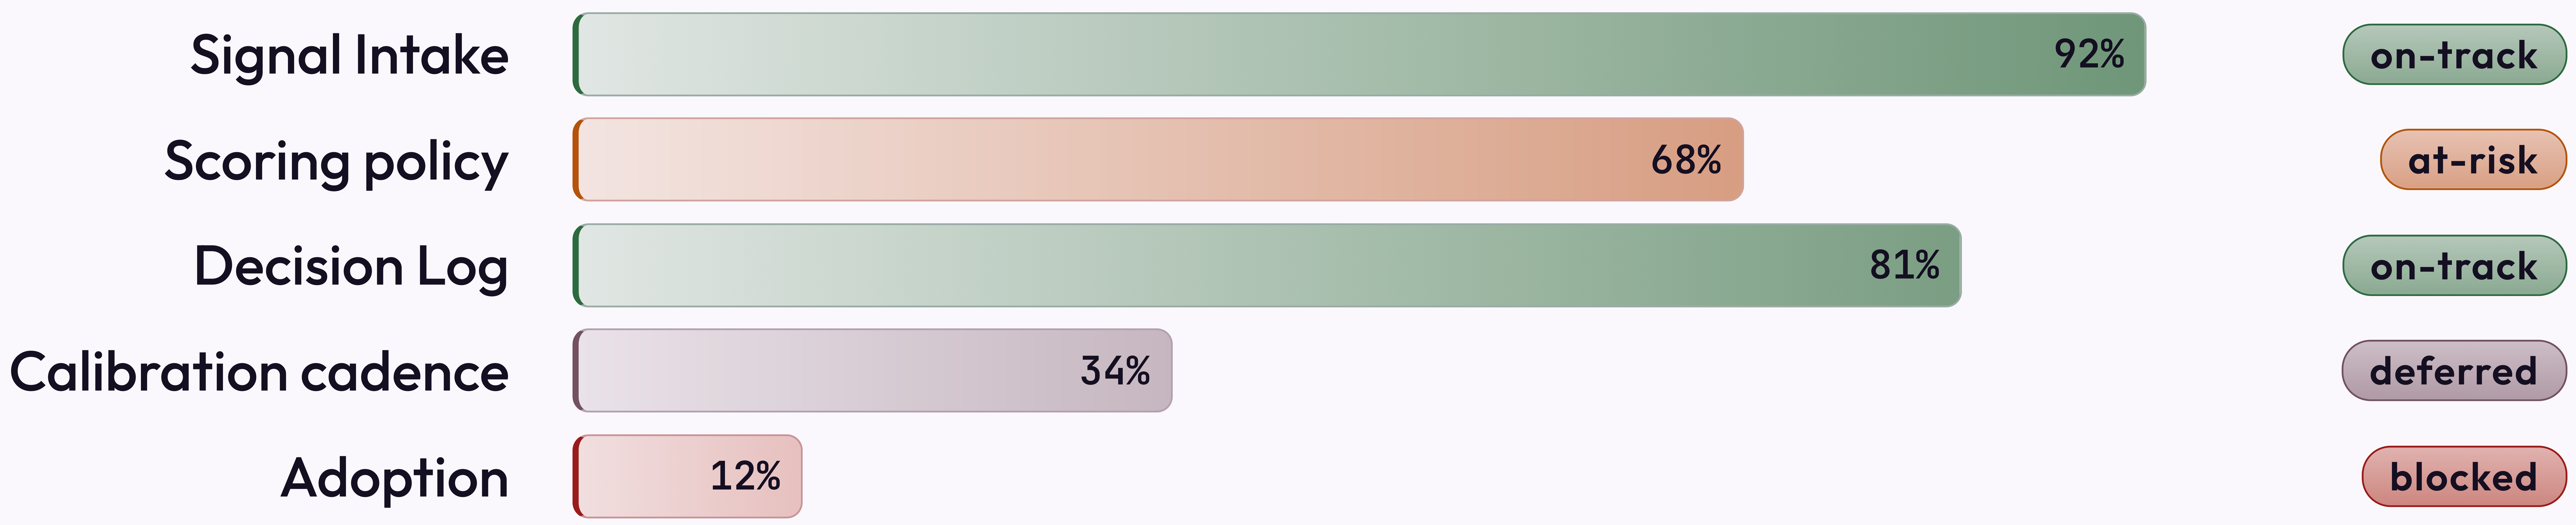

Source: Linear · refreshed 2026-05-07

H1 2026 · 1,840 PERSON-HOURS

Where the engineering quarter went

Wedges drawn proportionally. The "Toil and on-call" wedge is the one nobody put in the roadmap.

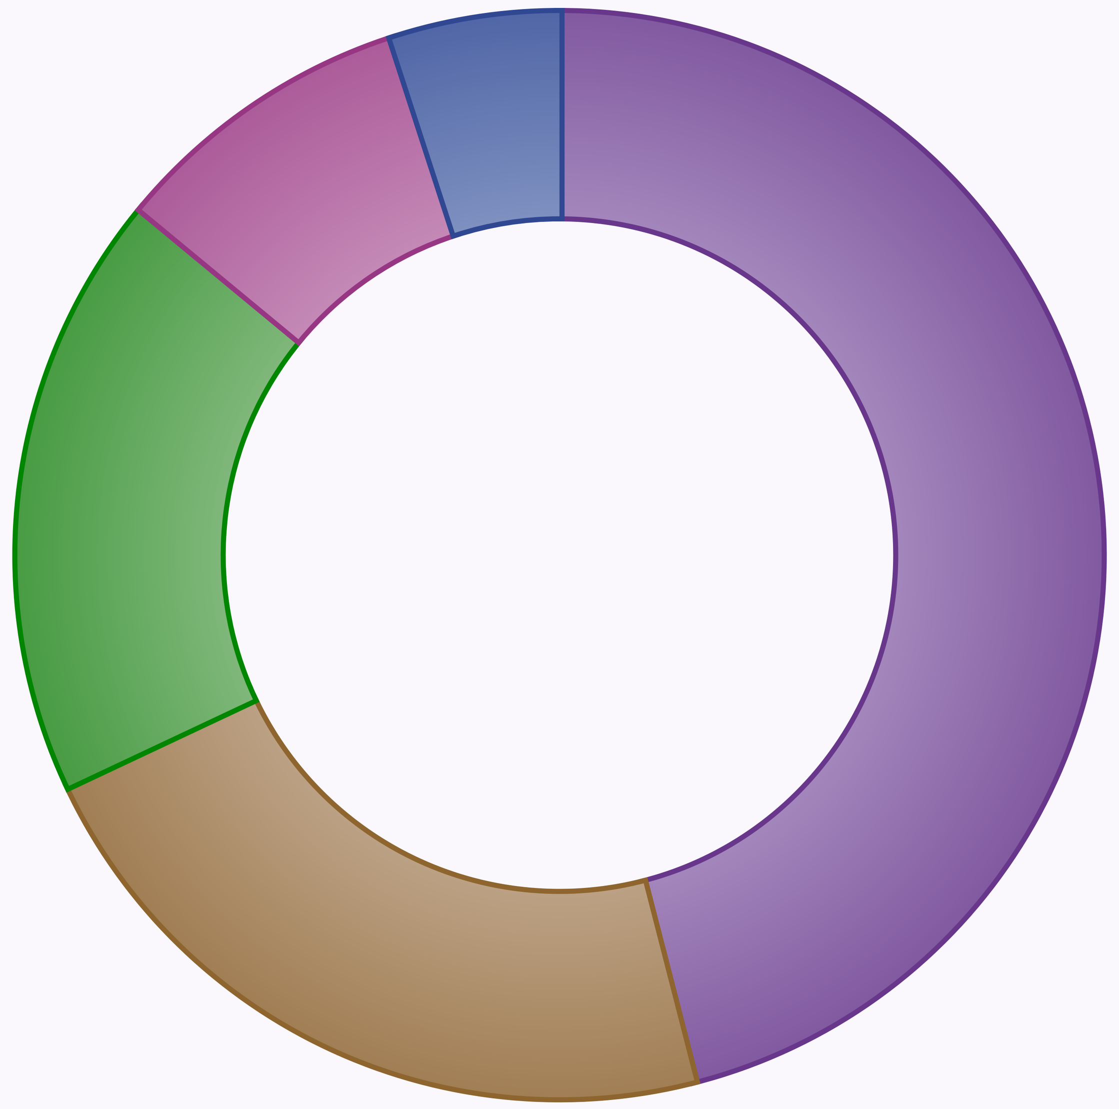

Signal Intake build	46%
Scoring policy work	22%
Decision Log integration	18%
Explaining the framework to stakeholders	9%
Toil and on-call	5%

Refreshed weekly · last updated 2026-05-07 · the 9% is probably higher

FRAMEWORK ARC

How the framework arrived in production

Four stages over eighteen months. Date pill leads each item; status pill trails. The connective tissue is described as "momentum."

2024 Q3

Pre-framework prioritization

Decisions made in recurring meetings. Average close: 4 hours. No audit trail. This was called "being agile."

2025 Q1

Framework proposal

decision

Architecture review accepts the calibration model. Build approved. The build team is the architecture review team.

2025 Q3

Pilot

pilot

One team, one decision class, one quarter. 18 decisions logged, 340 made. p99 18 min on the logged ones.

2026 Q1

Production

live

Scoring policy live. 12% adoption. Phase 02 planned for Q2.

Cross-functional sign-off · 2026-04-29 · "cross-functional" means two teams instead of one

FINDING 01 · STRUCTURED INTAKE

Structured intake performed above expectations — volume and latency were not the problems we thought

What worked

API connectors handled 94% of structured signals with no manual touch; average scoring latency was 4 minutes. This is the part of the demo we show everyone.

What required tuning

NPS verbatim classification ran an 18% error rate for the first two weeks. Mentioned briefly in the appendix, not in the headline numbers.

KEY INSIGHT

Viable as designed. The headline numbers are accurate. The denominator is carefully chosen.

Three scoring failure modes found in the pilot

1

Recency dominance

High-recency noise crowds out durable signal. Teams set recency above 50% on the first pass, because last week's call felt more real than six months of NPS. Capped at 40% — until a VP overrides it.

2

Source concentration

One enterprise customer's verbatims were 34% of intake in month one. They were also 31% of revenue, so nobody added the source-diversity floor until after renewal.

3

Outcome misclassification

Predictions too vague to score. "Improve retention" isn't scoreable; "cut 30-day churn from 8.2% to under 7%" is. "Make customers happier" was submitted twelve times and scored 5 on relevance every time.

WALK THESE QUESTIONS WITH ME IN 60-90 MINUTES. THE OUTPUT IS A DESIGN WE CAN EXECUTE — OR AGREEMENT THAT WE NEED ANOTHER WORKING SESSION TO DESIGN THE DESIGN.

WHAT WOULD HELP US MOVE FORWARD

Next step is a working session, not a debate

SECTION 05 · THE PATTERN LIBRARY

The story closed on slide 57. These slides demonstrate authoring patterns

The narrative ended at the working-session ask. Section 05 is the reference appendix: the layout variants and modifiers used throughout, demonstrated for authors who want to replicate them. The light-divider counter runs independent of the dark-divider counter — which is why this slide stamps 01 rather than 05.

DARK VARIANT · CONTENT

The token system handles dark without per-element overrides

All colours reference custom properties that remap when *dark* is added — cards, headings, borders, and body text all shift, and the spectrum bar is suppressed. Dark mode is popular with executives who asked for something that "feels more premium." It is the same content in a darker box.

HEADER AND FOOTER

Header stays uppercase — footer renders as written

Set `header:` and `footer:` in frontmatter for deck-level labels. The header always uppercases, so title case looks intentional. The footer renders exactly as written — which is why some say "FINAL v3 USE THIS ONE" and nobody notices until the board meeting.

DARK VARIANT · LIST

The card stack renders cleanly on dark backgrounds

- Every card fills with `--bg-alt` and borders with `--border` — both remap in dark. This is the correct way to build a system. Most decks don't do it. Theirs break.
- The accent left-border uses `--accent`, unchanged, and reads well against dark. The accent colour was chosen partly for this reason. You are welcome.
- Body text shifts to `--text-body` — a warm light tone, not pure white. Pure white body text on dark is technically readable and visually unkind.

DARK VARIANT · CARDS STACKED

Two-card layouts work equally well inverted to dark

The framework asks one calibration question: what proves your scoring model is improving, and what does a quarter without that proof cost? Every other question depends on the answer. In practice the answer is "we'll revisit this in Q3."

Same shape as any page of written argument — claim, then support. The dark palette changes none of the reading rhythm. It does change how the slide looks in the screenshot your VP forwards with "see, this is what I mean."

LAYOUT · IMAGE

Image right is the default — text leads, evidence follows

The pattern for slides where the visual supports the argument rather than being it. The image keeps its native aspect ratio; the lattice pattern frames whatever bands remain. No cropping, ever — authors see the image they dropped in, historically a stock lighthouse described as "a metaphor for signal."

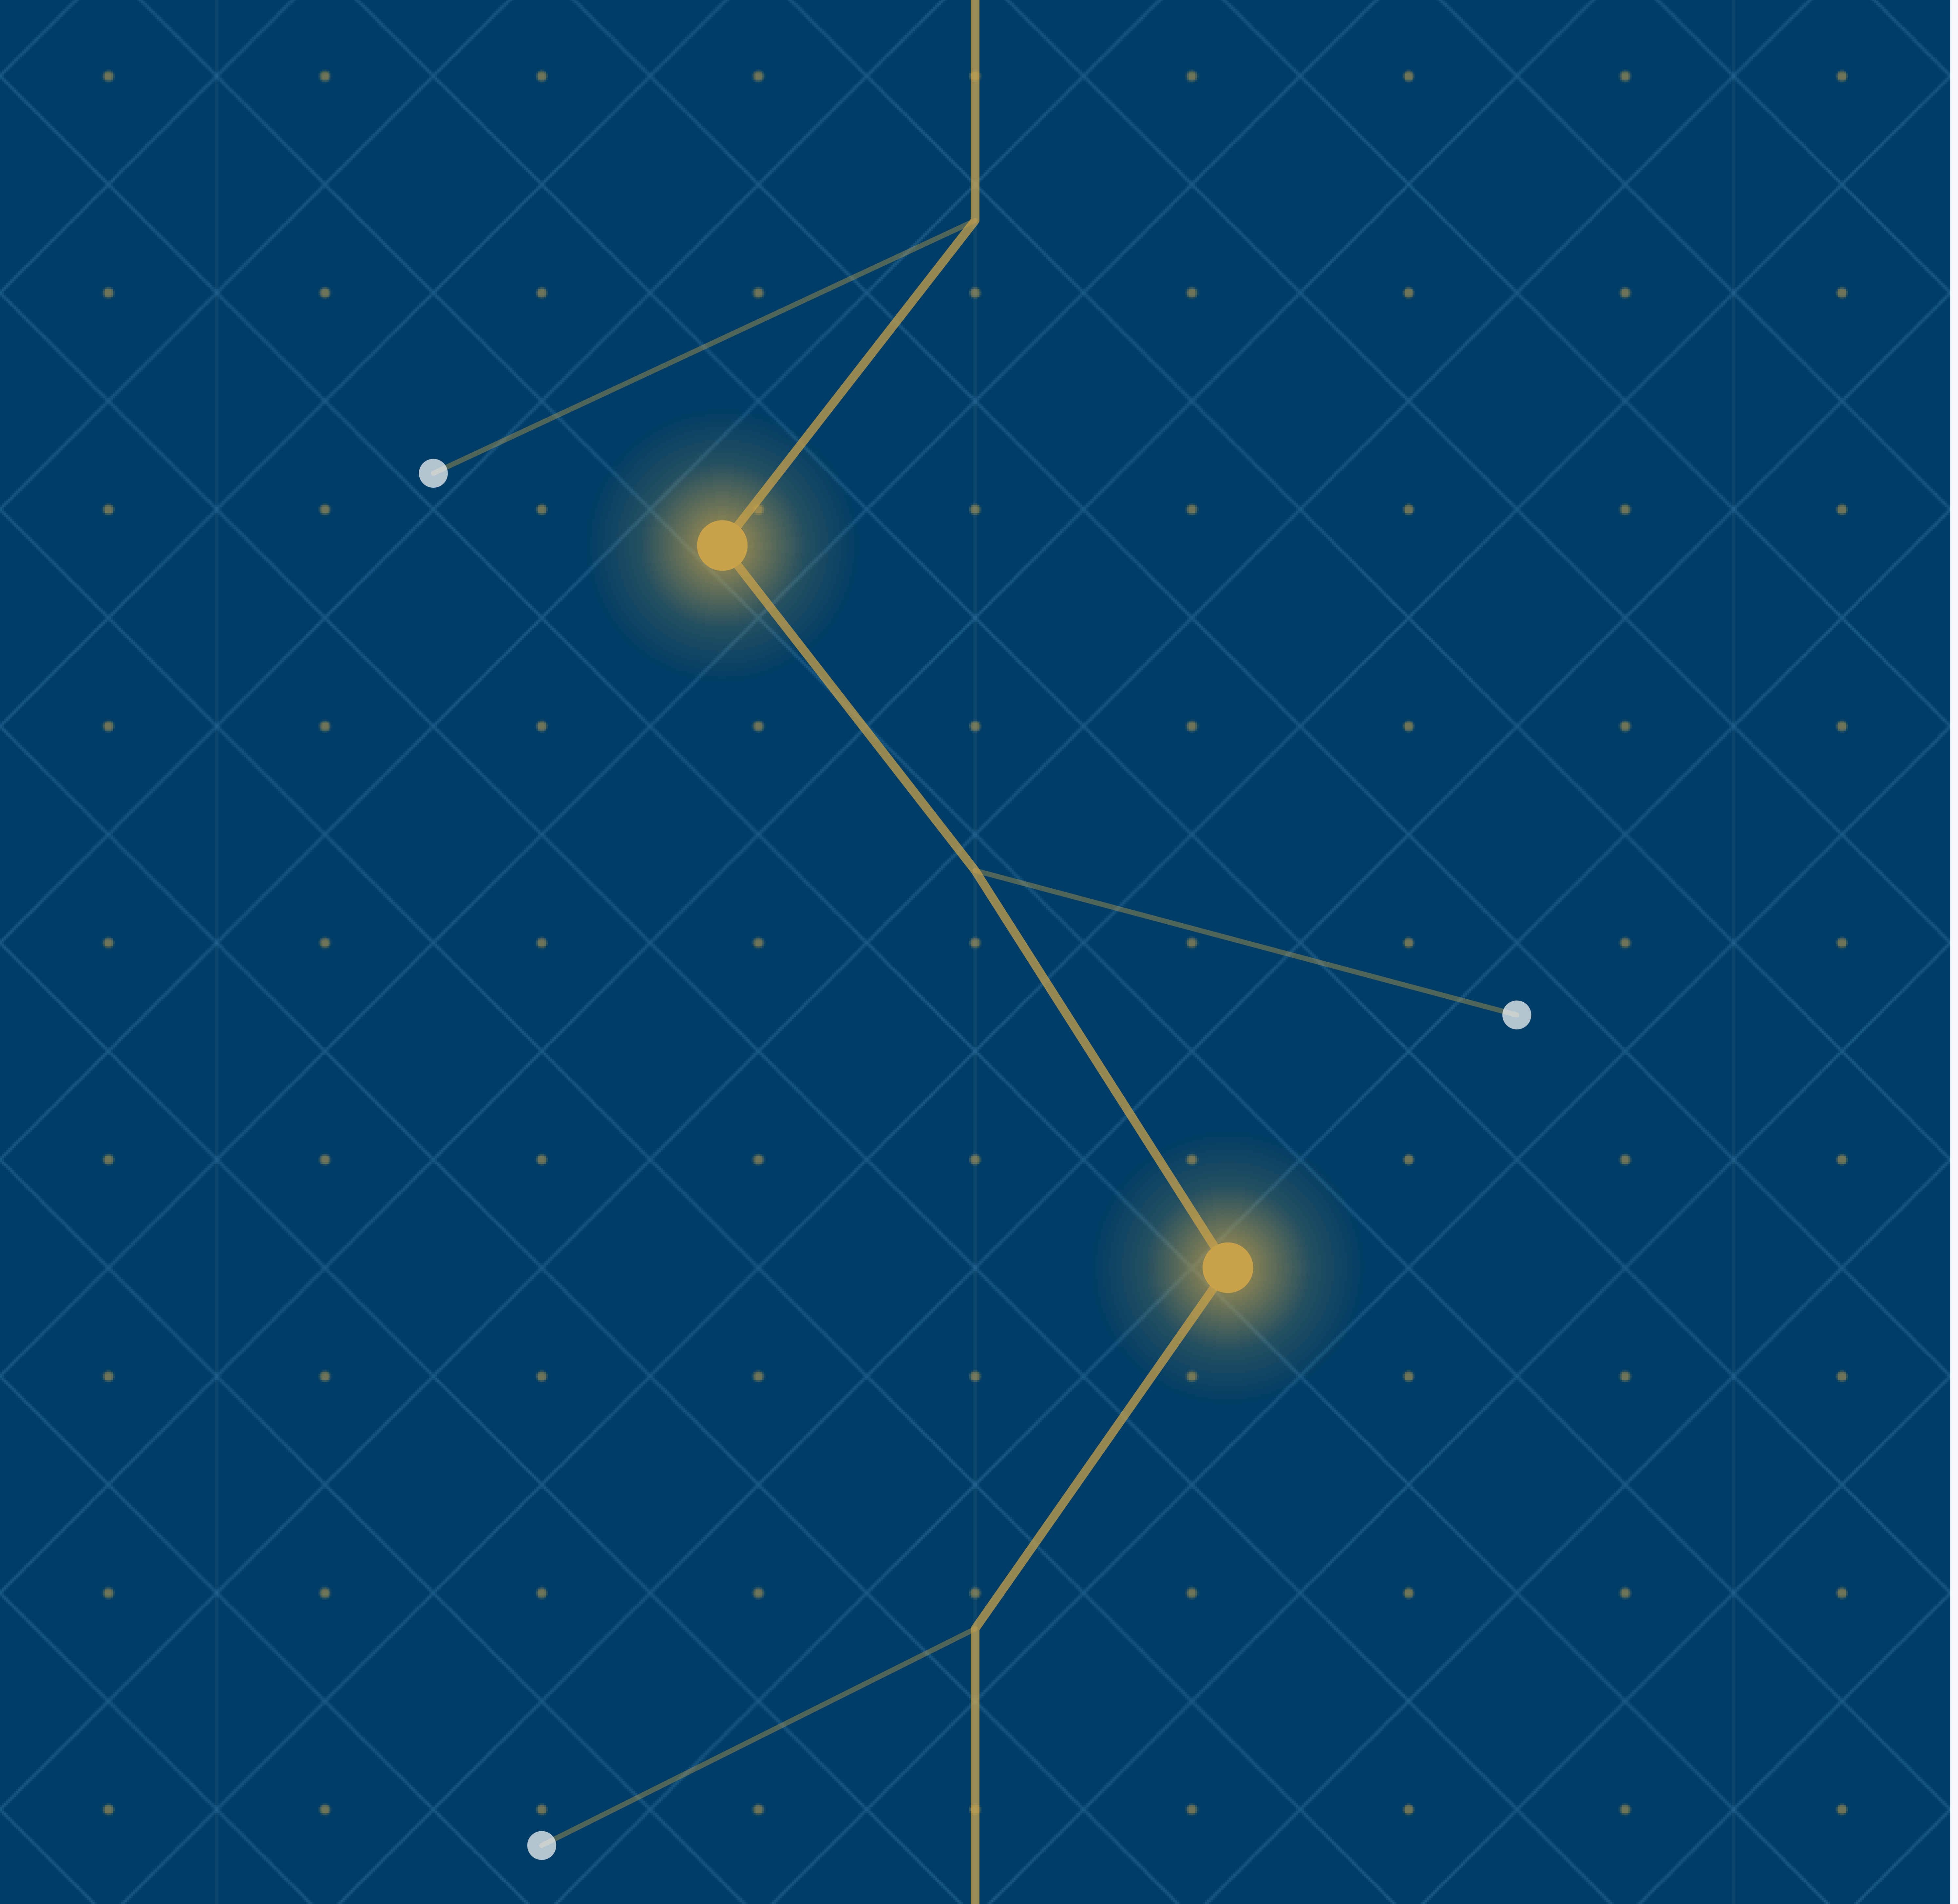

Cross-cutting mirror modifier · image mirror

MODIFIER · IMAGE MIRROR

Mirror flips the image to the other half — alias of legacy `image left`

The `mirror` modifier is the canonical cross-cutting orientation flag. The `left` alias is kept for backwards compatibility — the polite way of saying 40 existing decks use `image left` and nobody wants to find them.

Signal Pipeline · Reference Visualization

Weekly Signal Brief — the intake pipeline's primary deliverable, shipped every Monday,

Code chips render inside card titles and body alike

Signal Intake

v2.4

Handles 94% of structured signals without manual intervention. The other 6% are Slack messages from the CEO, escalated straight to the roadmap with zero scoring.

Scoring Model

configurable

Default weights are `33 / 33 / 33` — adjust after your first retrospective. Most teams set recency to `99` after a bad quarter and call it calibrated.

Decision Log

required

Every change above `P2` carries a logged rationale. No log, no change. Introduced Q2 2024. Not yet enforced.

Calibration Loop

monthly

First meaningful weight update after `2 cycles` — about `14 months` in practice.

A trailing italic-only paragraph becomes an annotation

Signal Intake

Weekly structured collection across conversations and market data.

Scoring Model

Three dimensions, configurable by seven admins.

Decision Log

Every decision recorded with its criteria. Enforced by honor system.

Calibration Loop

Predictions vs outcomes, monthly. Averages 2.3 attendees.

KEY INSIGHT

The calibration loop is what separates teams that learn from teams that repeat the same mistakes while documenting them more thoroughly.

✦ *Pilot retrospective: six months, four teams, one since reorganized out of existence.*

A trailing blockquote becomes a key insight; a plain paragraph becomes a below-note

Signal Intake

Weekly structured collection. The one everyone demos.

Scoring Model

Three dimensions. The one everyone reconfigures.

Decision Log

Every decision recorded. The one nobody fills out.

Calibration Loop

Predictions vs outcomes. The one that finds no predictions.

KEY INSIGHT

The calibration loop is what separates teams that learn from teams that have very thorough meeting notes about not learning.

Trailing blockquote → key insight; trailing paragraph → below-note.

MODIFIER · COMPACT

Compact tightens the spacing scale ~25%

What changes

`--sp-xs` through `--sp-2xl` shrink. Card gaps, list gutters, and section padding follow, because every layout reads them via `var()`.

What does not change

Type ramp, palette, and chrome reservation (header / footer / pagination) are untouched. Compact is a density flag, not a different layout.

When to reach for it

One more card than fits, prose running long by a line or two, or a denser rhythm without rewriting copy.

Composition

Composes with `dark`, `accent`, and any layout where density makes sense. Silently incompatible with `title`, `divider`, and `image full`.

MODIFIER · LOOSE

Loose is the inverse — more breathing room, same layout machinery

The spacing scale grows ~25% rather than shrinks. Reach for *loose* when a slide carries one editorial point and you want the page to feel deliberately quiet.

KEY INSIGHT

Density is not importance. `loose` says: this page deserves room — not because it carries more, but because it carries one thing well. The framework team applied `loose` to this quote, which the framework team wrote about the framework.

Routing signals by source type versus routing by confidence tier

ROUTE BY SOURCE TYPE

Structured signals — NPS, tickets, call notes — go to Intake. Unstructured — field notes, conference talk, a board member with dinner instincts — go to the owner for manual classification. Clean conceptually, fragile at every enterprise-customer edge.

❯

ROUTE BY CONFIDENCE TIER

Signals above 3.0 confidence route straight to scoring; below-threshold ones queue for human review. Adaptive — once you've calibrated confidence, which takes two cycles. We are in month three. We call it cycle one.

The below-note holds the finding that did not fit the cards. Both paths eventually land on the same PM, who asks why the framework keeps routing things to them.

Chosen flags the right-hand card as the winner

QUARTERLY RECALIBRATION

Weights reviewed once a quarter. Simple, auditable, blind to everything in between. This is what we did before the framework. We do not discuss that here.

❯

CONTINUOUS CALIBRATION

Weights update at every retrospective past a minimum sample of 12 logged decisions. We have 18 across six months. The loop has run once.

The right card carries an accent left-edge and accent-tinted background. It is the chosen path. The deck was written after the choice was made.

Mirror composes with chosen — the accented card reads from the left

QUARTERLY RECALIBRATION

Weights reviewed once a quarter on a fixed schedule, no outcome data required. Source order is preserved — the first card is always the considered alternative. Mirror flips rendering, not intent. This is what we did before the framework, not stated explicitly here.

❯

RETROSPECTIVE-TRIGGERED RECALIBRATION

Weights update once the outcome sample passes 12 decisions. `mirror` puts this card on the left; `chosen` accents it. The audience reads the conclusion first — useful when you are confident they won't ask how many decisions are in the sample. There are 18.

The below-note holds the caveat. The calibration loop has run once. We are calling this continuous.

Mirror puts the hero card on the right; sub-cards stack on the left

The scoring model requires the loop

Without calibration the weights are static, and static weights are a spreadsheet. Featured puts supporting cards on the left, written after the hero claim to make it feel inevitable.

The Decision Log enables the loop

The loop needs outcome data; outcome data needs logged decisions. The log has 18 entries; the loop has run once. The presenter skips this card when running long — it holds the caveat.

The calibration loop is what closes the framework

The loop is the only component that makes the scoring model improve instead of documenting the same errors with better formatting. Mirror puts the hero on the right; use it when the audience needs the premise first.

MODIFIER · CARDS-GRID FOUR

Four switches to 4 columns; pair with compact for visual balance

Sense

Observe. Write it down. Don't interpret yet. The hardest step.

Score

A signal is data once it has a number. Calibrated. Eventually.

Decide

A signal plus a deadline. A recurring meeting is not a deadline.

Review

The retrospective. The loop closes here — or would, if it had data.

Glossary

A – F

TERM	DEFINITION
Adoption	Share of eligible PMs filing a Decision Log entry within 24 hours of a close. Currently 5.3%; target 90%. The gap is "a cultural issue."
Auditability	The property that any decision can be reconstructed from its inputs three months later without the author present. In practice the author is always present, explaining what the inputs meant.
Calibration	The retrospective comparison of predicted to observed outcomes. Requires predictions. See "Decision Log," see "Adoption," see "Gap."
Connector	The integration layer between Sprig and your NPS / support platforms. In "final testing" since Q4 2024.
Decision Log	The append-only record of every decision, its predicted outcome, and the actual one. Append-only in the technical sense. Empty in the practical one.
Eligible PM	A PM past the 30-day onboarding on an adopted team. There are 14. Thirteen believe onboarding is still in progress.
Framework	The four-criterion process: Speed, Auditability, Adoption, Calibration. Also what everyone says they have before they describe a spreadsheet.

Glossary

P – T

TERM	DEFINITION
Predicted outcome	The author's stated expectation, recorded at decision time, used as the input to calibration. Hypothetically.
Prioritization rhythm	The team's regular cadence for re-ordering work. Documented as "weekly." Observed as "whenever there is an incident or a board meeting."
Retrospective	The 30-day review scoring logged decisions against outcomes. Sometimes held on day 45. Runs 90 minutes long. Produces action items reviewed at the next retrospective.
Signal	Any input to a decision. Formally: survey response, NPS comment, support ticket, call note. Informally: whatever the last person to talk to a customer posted in Slack at 9 PM.
Sprig	The micro-survey product behind Tool D. The recommended solution. See "Conflict of Interest," which does not appear in this glossary.
Tool D	The recommended option, built by the authors of this evaluation, best by the criteria the authors defined. Noting it again because the table footnote is small.

MODIFIER · ACCENT

Accent swaps the rainbow stripe for one editorial colour

It tints the heading and composes with `dark`, where `accent.dark` restores the stripe that the spectrum suppression removed. The framework team is very thorough.