FIT LADDER · SPLIT WHAT CAN'T PAGINATE

The read-across carousel

A comparison can't be cut down the middle — so when it won't fit, it becomes a sequence.

What the carousel solves

- A list overflows? Split it between items — that is **auto-split**
- A **comparison** reads *across* its two sides — slice it and the meaning is gone
- So the engine **re-authors** it as a sequence: a cover, one page per side, a verdict
- Each frame stands alone, reads in order — no setting, no shrinking

A COMPARISON

Scoring model: before and after the calibration loop

Two readings →

SCORING MODEL: BEFORE AND AFTER THE CALIBRATION LOOP

READING ONE

Before Calibration

Equal weights — confidence, recency, and relevance each contribute 33%. At least it is honest that we are guessing. No team has changed them in eighteen months.

SCORING MODEL: BEFORE AND AFTER THE CALIBRATION LOOP

READING TWO

After Calibration

Weights reflect each team's historical accuracy — but the team keeps changing, and the history is one atypical quarter. So most teams keep the 33s and call them calibrated.

THE SHIFT

“The shift takes two retrospective cycles — 60 days from adoption, or 14 months, depending on who you ask.”

A COMPARISON

The evaluation came down to a question the vendors could not answer

Two readings →

THE EVALUATION CAME DOWN TO A QUESTION THE VENDORS COULD NOT ANSWER

READING ONE

Buy a vendor framework

Three vendors evaluated. None expose calibration weights to the customer — the criterion that eliminated all three. It was added to the rubric after the team chose to build.

THE EVALUATION CAME DOWN TO A QUESTION THE VENDORS COULD NOT ANSWER

READING TWO

Build the framework in-house

Owns the scoring policy, the calibration loop, and the timeline. The window closes in eighteen months; a vendor cutover would eat nine of them.

THE SHIFT

"The conclusion came first; the evaluation was built to hold it. Everyone in the room knew it."

Decisions used to require a quarterly re-litigation

BEFORE

Decisions lived in the room they were made in. Six months on, nobody could say why a project was killed — only that someone senior had felt strongly.

↓

AFTER

Every decision is logged with its signals, its options, and the bet it made. Teams still relitigate, but now there is a record — so the argument is shorter.

SECTION 01 · DEEP DIVE

Scoring Model Deep Dive

The most configurable component — a feature or a warning sign, depending on your team. Six dimensions:

SCORING MODEL DEEP DIVE

Confidence

Independent corroborating sources, 1–5. Enterprise customers count as 1, regardless of volume.

Recency

Time-decay, configurable half-life. Set to two weeks, then surprised only recent news scores.

Strategic Relevance

Owner-scored, 1–5. A 5 correlates with whoever is presenting the roadmap this quarter.

SCORING MODEL DEEP DIVE

Source Diversity

Rewards breadth. In practice rewards whoever subscribes to the most newsletters.

Volatility

Penalizes signals that swing. Also penalizes the only signal that caught the last surprise.

Owner Bias

The correction nobody applies, scored by the owner it is meant to correct.

Two shapes, two treatments

- A **comparison** (compare-prose) → a cover, a page per side, a verdict
- A **feature panel** (split-panel) → the feature gets a cover, its points flow on
- Both stage instead of slicing; both opt in with `autosplit: on`
- The layout *owns* its split-forms — declared in its manifest

More slides, never a clipped comparison

Read-across splits by staging, not slicing