

BLAMELESS RETRO • 9 JUNE 2026

The 2 June payments outage

What happened, what we learned, and the four changes we are making.

On 2 June, payments were down for 47 minutes — here is what we learned.

A configuration change took the payments service offline during the evening peak. We recovered in 47 minutes. This retro is blameless: the goal is to understand the system that let it happen and fix it, not to find a person to blame.

The blameless frame: we fix systems, not people.


Good people made reasonable decisions with the information they had.

The outage happened because our system let a single unreviewed change reach production at the worst possible time. We focus on the guardrails that were missing, not the hands on the keyboard.

- Every action item targets a process or a safeguard, never an individual.
- The engineer who made the change helped write this retro.

2 JUNE 2026 • ALL TIMES LOCAL

How the incident unfolded, minute by minute.



CUSTOMER IMPACT

47 min

of failed checkouts during evening peak —
roughly 12,000 transactions affected, all
later recoverable.

What went well, and what went poorly.

Went well • Detection.

Customer reports reached us within two minutes

Went well • Response.

War room and rollback ran cleanly once declared

Went poorly • Prevention.

The change skipped staging entirely

Went poorly • Alerting.

We learned from customers before our own monitors

Root cause: a manual change path that bypassed every safeguard.

The routing rule was changed through a console that writes straight to production. There was no required review, no staging step, and no alert on the resulting error spike — so three independent safeguards were all absent on the same path.

- The console predates our deployment pipeline and was never retired.
- Error-rate alerting covered the API but not the routing layer.

RESPONSE • THIS INCIDENT VS OUR TARGETS

How our response compared to our incident targets.

TARGETS FROM THE INCIDENT-RESPONSE RUNBOOK, LAST REVISED Q1 2026.

9 min	38 min	47 min	100%
TO DECLARE (TARGET 5)	TO MITIGATE (TARGET 30)	TOTAL DOWNTIME	TRANSACTIONS RECOVERED

The four changes we are making, with owners.

STEP 01

Retire the console

Direct production writes are disabled this week; all changes route through the pipeline.
Owner: SRE.

STEP 02

Require staging

No config change reaches production without a staging deploy and a second reviewer.
Owner: Platform.

STEP 03

Alert the routing layer

Error-rate alerts now cover routing, paging on-call within 60 seconds. Owner: SRE.

STEP 04

Rehearse rollback

A monthly game-day drills the rollback path so it stays fast. Owner: Eng leads.

Action-item status as of this retro.

☒	Incident timeline reconstructed	complete
☒	Customer comms sent	complete
☐	Direct-write console disabled	in progress
☐	Staging gate enforced in pipeline	this sprint
☐	Routing-layer alerts live	this sprint
☐	First rollback game-day scheduled	July

INCIDENT REVIEW · SRE@EXAMPLE.COM

**We fix the system so
this can't happen the
same way twice.**