

BLAMELESS RETRO • 9 JUNE 2026

The 2 June payments outage

What happened, what we learned, and the four changes we are making.

On 2 June, payments were down for 47 minutes — here is what we learned.

A configuration change took the payments service offline during the evening peak. We recovered in 47 minutes. This retro is blameless: the goal is to understand the system that let it happen and fix it, not to find a person to blame.

The blameless frame: we fix systems, not people.


Good people made reasonable decisions with the information they had.

The outage happened because our system let a single unreviewed change reach production at the worst possible time. We focus on the guardrails that were missing, not the hands on the keyboard.

- Every action item targets a process or a safeguard, never an individual.
- The engineer who made the change helped write this retro.

2 JUNE 2026 • ALL TIMES LOCAL

How the incident unfolded, minute by minute.



CUSTOMER IMPACT

47 min

of failed checkouts during evening peak —
roughly 12,000 transactions affected, all
later recoverable.

What went well, and what went poorly.

Went well • Detection.

Customer reports reached us within two minutes

Went well • Response.

War room and rollback ran cleanly once declared

Went poorly • Prevention.

The change skipped staging entirely

Went poorly • Alerting.

We learned from customers before our own monitors

Root cause: a manual change path that bypassed every safeguard.

The routing rule was changed through a console that writes straight to production. There was no required review, no staging step, and no alert on the resulting error spike — so three independent safeguards were all absent on the same path.

- The console predates our deployment pipeline and was never retired.
- Error-rate alerting covered the API but not the routing layer.

RESPONSE • THIS INCIDENT VS OUR TARGETS

How our response compared to our incident targets.

TARGETS FROM THE INCIDENT-RESPONSE RUNBOOK, LAST REVISED Q1 2026.

9 min	38 min	47 min	100%
TO DECLARE (TARGET 5)	TO MITIGATE (TARGET 30)	TOTAL DOWNTIME	TRANSACTIONS RECOVERED

The four changes we are making, with owners.

STEP 01

**Retire the console —
Direct production
writes are disabled this
week; all changes route
through the pipeline.**

Owner: SRE

STEP 02

**Require staging — No
config change reaches
production without a
staging deploy and a
second reviewer.**

Owner: Platform

STEP 03

**Alert the routing layer
— Error-rate alerts now
cover routing, paging
on-call within 60
seconds.**

Owner: SRE

STEP 04

**Rehearse rollback — A
monthly game-day
drills the rollback path
so it stays fast.**

Owner: Eng Leads

Action-item status as of this retro.

✓	Incident timeline reconstructed	complete
✓	Customer comms sent	complete
—	Direct-write console disabled	in progress
○	Staging gate enforced in pipeline	this sprint
○	Routing-layer alerts live	this sprint
○	First rollback game-day scheduled	July

INCIDENT REVIEW · SRE@EXAMPLE.COM

**We fix the system so
this can't happen the
same way twice.**