

EVIDENCE • CANVAS • SERIES


progress

Horizontal progress bars — one row per item, percentage filled.


H1 2026 · PHASE 1 READINESS

Phase 1 readiness, by workstream.

Snapshot at 14:00 UTC. Status pills reflect the most optimistic reading of the available data.




Stress test — eight workstreams, full status range, one honest pill.



H1 2026 · PHASE 1 READINESS

Phase 1 readiness, by workstream.


Snapshot at 14:00 UTC. Status pills reflect the most optimistic reading of the available data.



H1 2026 · PHASE 1 READINESS

Phase 1 readiness, by workstream.


Snapshot at 14:00 UTC. Status pills reflect the most optimistic reading of the available data.



H1 2026 · PHASE 1 READINESS

Phase 1 readiness, by workstream.

Snapshot at 14:00 UTC. Status pills reflect the most optimistic reading of the available data.



When NOT to reach for progress.

Comparing unrelated metrics. Revenue % of target, latency vs SLO, and headcount fill aren't comparable on a shared bar scale. Use `kpi` for value/target/status tiles or `stats` for an independent metric row.

More than eight rows. Past eight workstreams the bars compress and the labels truncate. Split the view by owner or workstream group; the audience can't scan twelve bars at once anyway.

Decorative status pills. Don't invent new status words for tone. `on-track`, `at-risk`, `blocked`, `deferred`, `done` are the vocabulary the engine recognises; everything else renders as a plain pill and breaks the at-a-glance read.

RELATED COMPONENTS

See also.

`kpi` — value + target + status tiles, not a single percent

`stats` — independent headline metrics, no completion scale

`gantt` — the rows are time-bound and need a date axis

`checklist` — binary done / not-done across a flat list

`timeline-list` — the workstreams complete in sequence, not in parallel