

PRODUCT STRATEGY • Q3 2025

From Signal to Strategy

A decision framework for product leaders navigating market uncertainty

SECTION 01 • FOUNDATIONS

The landscape has shifted.
Here is what that means
for us.

MODULE 02

Before we score signals, we need to agree on what a signal is.


The word is overloaded. We use it to mean anything from a customer complaint to a macro trend. This framework requires a tighter definition.

The window for differentiation is narrowing.

Three converging forces — commoditized infrastructure, compressed release cycles, and rising customer switching costs — have reduced the average durable advantage window from 36 months to under 14. Teams that cannot identify signal from noise in that window will consistently miss timing.

How signals move from input to decision.

Four-stage processing pipeline – weekly cadence



Six months of results across four product teams.

MEASURED AGAINST PRE-FRAMEWORK BASELINE, SAME TEAMS, SAME MARKET CONDITIONS.

73%
FASTER CLOSE

4.2×
SIGNAL RECALL

18
DECISIONS LOGGED

91%
TEAM ALIGNMENT

The framework has four components.

Signal Intake

Weekly structured collection across customer conversations, market data, and competitive moves. Normalized into a common schema before scoring.

Scoring Model

Each signal scored on three dimensions: confidence, recency, and strategic relevance. Weights are team-configurable and reviewed quarterly.

Decision Log

Every decision recorded with the signals that informed it, the options considered, and the criteria applied. Feeds the calibration loop.

Calibration Loop

Monthly retrospective that compares predicted outcomes to actual outcomes and adjusts scoring weights accordingly.

Code in card headers and body text.

Signal Intake

v2.4

Handles 94% of structured signals without manual intervention. Average latency: 4 min from ingestion to scored entry.

Scoring Model

configurable

Three dimensions: confidence, recency, relevance. Default weights are 33 / 33 / 33 — adjust after your first retrospective.

Decision Log

required

Every prioritization change above P2 must carry a logged rationale. No log, no change.

Calibration Loop

monthly

Compares predicted outcomes to actuals. First meaningful weight update happens after 2 cycles.

Signal Intake produces three outputs.

1

Weekly Signal Brief

A ranked list of the top 10 signals from the prior week, with confidence scores and source attribution. Distributed to product leads every Monday morning.

2

Anomaly Alerts

Real-time flags when a signal exceeds the 2σ threshold on any dimension. Routed directly to the accountable PM with a 4-hour response SLA.

3

Monthly Signal Index

The source of truth for the calibration loop. A complete record of all signals logged, scored, and resolved in the prior month. Required reading before each retrospective.

Two failure modes the framework is designed to prevent.

False signal amplification

A single loud voice — one enterprise customer, one analyst report, one competitive announcement — dominates the decision without being weighed against the full signal set. The scoring model prevents any single source from exceeding 30% of the total signal weight in a given decision.

Signal hoarding

Teams collect signals but do not log decisions, so the calibration loop has nothing to learn from. The Decision Log is a required artifact for any prioritization change above P2 severity. No log, no change.

Two intake modes for different signal types.

Structured Intake

Signals with clear schema: NPS verbatims, support ticket categories, feature request volumes, win/loss notes. Ingested automatically via API connectors. Scored on arrival. Zero manual handling.

Unstructured Intake

Signals without schema: field observations, conference conversations, analyst briefings, competitive demos. Require human classification before scoring. Routed to the signal owner for a 48-hour classification window.

Scoring model: before and after the calibration loop.



The shift from equal weights to calibrated weights takes two retrospective cycles — roughly 60 days from adoption.


“

The signal was always there. We just didn't have a system that forced us to look at it before we'd already decided.

”

— Head of Product, Pilot Team 3

How a decision moves through the framework.



What the framework does not do.

It does not make decisions — it structures the information that humans use to decide.

It does not replace customer discovery — it scores and routes what discovery surfaces.

It does not work without the Decision Log — calibration requires outcome data to learn from.

It does not guarantee alignment — it surfaces disagreement earlier, which still requires resolution.

It does not scale down to individual feature decisions — it is designed for prioritization above P2.

Four things that must be true before you begin.

- 01 You have a regular prioritization cadence — at minimum monthly.
- 02 At least one person owns signal collection full-time or as a primary responsibility.
- 03 Leadership has agreed to log decisions with rationale, not just outcomes.
- 04 You have 90 minutes per week to run the intake and scoring process.

CALIBRATION RESULT • 6-MONTH PILOT

14X

Return on signal investment — measured as decisions that reached the right outcome on the first attempt, versus the baseline rate before the framework was adopted.



Scoring Model Deep Dive

What this section covers

The scoring model is the most configurable component. This section covers the three dimensions, how weights are set initially, and how calibration updates them over time.

- 1

Confidence
How many independent sources corroborate the signal. Ranges 1–5.
- 2

Recency
Time-decay applied from signal date to scoring date. Half-life is team-configurable.
- 3

Strategic Relevance
Manual score from the signal owner. Ranges 1–5. Requires justification above 4.

WALK THESE QUESTIONS WITH ME IN 60-90 MINUTES. THE OUTPUT IS EITHER A DESIGN WE CAN EXECUTE, OR A SHARED LIST
OF WHAT NEEDS MORE WORK BEFORE WE COMMIT.

WHAT WOULD HELP US MOVE FORWARD

Next step is a working
session, not a debate.

FINDING 01 • STRUCTURED INTAKE

Structured intake performed above expectations — volume and latency were not concerns.

What worked

API connectors handled 94% of structured signals without manual intervention. Average scoring latency was 4 minutes from ingestion. Schema normalization held across all five connected sources.

What required tuning

NPS verbatim classification had an 18% error rate in the first two weeks. Required a training pass on the classification model before accuracy reached the 92% target.

Viable as designed — NLP classification requires a 2-week warm-up period on new deployments.

Key insight works on any card-bearing layout.

Signal Intake

Weekly structured collection across customer conversations, market data, and competitive moves.

Scoring Model

Each signal scored on three dimensions: confidence, recency, and strategic relevance.

Decision Log

Every decision recorded with the signals that informed it and the criteria applied.

Calibration Loop

Monthly retrospective that compares predicted outcomes to actual outcomes.

The calibration loop is what separates teams that learn from teams that repeat the same mistakes.

Trailing blockquote becomes a key insight; trailing paragraph becomes a below-note with a hairline rule above it.

A trailing italic-only paragraph becomes an annotation.

Signal Intake

Weekly structured collection across customer conversations, market data, and competitive moves.

Scoring Model

Each signal scored on three dimensions: confidence, recency, and strategic relevance.

Decision Log

Every decision recorded with the signals that informed it and the criteria applied.

Calibration Loop

Monthly retrospective that compares predicted outcomes to actual outcomes.

The calibration loop is what separates teams that learn from teams that repeat the same mistakes.

Source: pilot retrospective, six months across four product teams.

Three scoring failure modes found in the pilot.

1

Recency dominance

High-recency noise crowding out durable signal. Teams set recency weight above 50% in the first calibration pass. Corrected by capping recency weight at 40% until two calibration cycles complete.

2

Source concentration

Single-customer signals inflating confidence scores. One enterprise customer's verbatims represented 34% of all structured intake in month one. Corrected by adding a source-diversity floor to the scoring model.

3

Outcome misclassification

PMs logging outcomes too vague to score. "Improve retention" is not scoreable. "Reduce 30-day churn from 8.2% to below 7%" is.

Four requirements every decision system must meet.

01 **Speed**

Decisions must close within the window they are relevant to. Systems that add latency consume the value they exist to protect.

02 **Auditability**

Every prioritization decision above a threshold must carry a traceable rationale. Required for alignment and compliance.

03 **Adoption**

If the team won't use it weekly, calibration never runs and the model never improves. Ninety minutes per PM is the ceiling.

04 **Calibration**

The system must improve over time. A static scoring model is a spreadsheet with extra steps.

We evaluated four intake tools against the criteria.

Tool A • Chorus

- ✔ Speed
- ✖ Auditability
- ✔ Adoption
- ✖ Calibration

Strong call recording and summarization. No decision logging or calibration loop. Requires separate tooling for everything downstream of intake.

Tool B • Productboard

- ✖ Speed
- ✔ Auditability
- ✔ Adoption
- ✖ Calibration

Solid intake and prioritization. Decision logging exists but is manual and rarely used. No calibration mechanism. Setup takes 3–4 weeks.

Tool C • Notion

- ✔ Speed
- ✔ Auditability
- ✖ Adoption
- ✖ Calibration

Flexible enough to build the full system. But building it takes 40+ hours and the result is fragile. Teams abandon maintenance after the first quarter.

Tool D • Sprig + Decision Log

- ✔ Speed
- ✔ Auditability
- ✔ Adoption
- ✔ Calibration

Meets all four criteria within the 90-minute weekly budget. Reaches production in the same week it is adopted. Recommended.

The four tools side by side.

CRITERION	CHORUS	PRODUCTBOARD	NOTION	SPRIG + LOG
Speed	✓	✗	✓	✓
Auditability	✗	✓	✓	✓
Adoption	✓	✓	✗	✓
Calibration	✗	✗	✗	✓
Setup time	1 day	3–4 weeks	40+ hours	Same day

Evaluated against the same four teams and the same 90-minute weekly budget constraint.

Glossary

A – F

TERM	DEFINITION
Adoption	Percentage of eligible PMs filing a Decision Log entry within 24 hours of a decision close.
Auditability	The property that any decision can be reconstructed from its inputs three months later without the original author present.
Calibration	The retrospective comparison of predicted to observed outcomes, used to score the framework's accuracy.
Connector	The integration layer between Sprig and your NPS / support platforms. Owns ingestion and tagging.
Decision Log	The append-only record of every prioritization decision, its predicted outcome, and the actual outcome at retrospective time.
Eligible PM	A PM whose team has adopted the framework and is past the 30-day onboarding period.
Framework	The four-criterion process: Speed, Auditability, Adoption, Calibration. The deck's central artifact.

Glossary

P – T

TERM	DEFINITION
Predicted outcome	The author's stated expectation, recorded at decision time, used as the input to calibration.
Prioritization rhythm	The team's regular cadence (weekly, biweekly) for revisiting and re-ordering work.
Retrospective	The 30-day review meeting where logged decisions are scored against observed outcomes.
Signal	Any qualitative or quantitative input to a decision — survey response, NPS comment, support ticket, sales call note.
Sprig	The micro-survey product used by Tool D to capture qualitative signal in-product.
Tool D	The recommended option in the four-tool comparison: Sprig combined with a lightweight Decision Log.

Two options with a connector and an explanatory note below.

OPTION A • LABEL

Body text describing the first option. Enough detail to fill the card naturally and show how the layout handles a few lines of prose.



OPTION B • LABEL

Body text describing the second option. The connector arrow between them implies direction or causality — before/after, input/output, cause/effect.

The below-note sits under the cards after a hairline rule. Use it for a single contextual sentence.

How to roll this out across your organization.

STEP 01

Pick one team and one decision type

Start with a team that already has a regular prioritization rhythm. Apply the framework only to a single decision category for the first 30 days.



STEP 02

Log everything, decide nothing differently

In the first month, do not change how you make decisions. Just log signals and decisions as you would have made them anyway.



STEP 03

Run your first retrospective

At day 30, score the logged decisions against outcomes. This is where the model gets its first calibration pass.



STEP 04

Expand to a second team

With one retrospective complete, you have evidence. Use it to onboard the second team with real data, not promises.

The six signal dimensions, what they measure, and how they are scored.

01	Confidence	Number of independent sources corroborating the signal	1-5 • AUTO-SCORED
02	Recency	Time-decay from signal date, configurable half-life	0.0-1.0 • AUTO-SCORED
03	Relevance	Alignment to current strategic bets, owner-scored	1-5 • MANUAL
04	Reach	Number of customers or segments affected	1-5 • AUTO-SCORED
05	Effort	Engineering and design cost to act on the signal	1-5 • MANUAL
06	Confidence delta	Change in confidence score since last scoring cycle	-5 TO +5 • AUTO

HEADER AND FOOTER

Header stays uppercase — footer renders as written.

Set `header:` and `footer:` in frontmatter for deck-level labels, or use per-slide comment directives. The header uses uppercase text-transform automatically, so you write it in any case. The footer renders exactly as written.

Wiring a signal into the framework is three lines of code; the onboarding is three months.

```
import { DecisionFramework } from "@company/signal-sdk";

const framework = new DecisionFramework({ configFile: "./framework.config.json" });

// Score a signal at intake
const score = await framework.score(signal, { dimensions: ["confidence", "recency", "relevance"] });

// Log every decision – calibration depends on it (nobody calls this in prod)
const entry = await framework.decisions.log(decision, { signals: [signal.id], rationale });
```


BEFORE & AFTER • SCORING MECHANICS

Spreadsheet-driven scoring versus framework-driven scoring.

BEFORE • THE HONEST SPREADSHEET

```
# Manual scoring. Auditable in the  
# sense that you can see who edited it  
import pandas as pd  
  
signals = pd.read_csv('./signals.csv')  
signals['score'] = signals.apply(  
    lambda r: 0.33*r.confidence + 0.33*r.recency + 0.33*r.relevance,  
    axis=1,  
)
```


AFTER • THE FRAMEWORK


```
# Calibrated weights, signed policy,  
# every score is audit-logged  
from decision_framework import Calibrator  
  
calibrator = Calibrator.load('./policy.json')  
for signal in calibrator.intake.unscored():  
    calibrator.score(signal)  
    calibrator.decisions.log_if_relevant(signal)
```


LAYOUT • IMAGE

Hand it any photo — the layout resolves the composition.

A moderate-aspect photo lands in the clean default: a floated card shaped to the asset, beside the argument. No modifier, no crop — the engine reads the asset's shape and the deck size and decides.





LAYOUT • IMAGE MIRROR

**Mirror flips the card —
image left, text right.**


`mirror` is the cross-cutting side modifier; `! [bg
left]` is preserved as a deprecated alias. The
composition is unchanged — only the side flips.



LAYOUT • IMAGE GALLERY

When a chart or screenshot must show in full, opt into `gallery`.

The whole asset is contained at native aspect on a clean `--bg-alt` matte with a placard caption — an editorial plate, not a placeholder. Use it for diagrams, schematics, and any asset where cropping would destroy meaning.




Signal Pipeline · Reference Visualization


A panorama matches the canvas, so it goes full-bleed — and the message rides a solid card, so it stays legible over any photo.

IMAGE LAYOUT • SPLIT

A tall photo keeps every pixel.

An extreme aspect would waste a card, so the layout shows it whole in a full-height column with the argument alongside — zero crop, the canvas filled.





Architecture dependencies — the editorial frame

statement rides the title on the photo over a scrim — a deliberate, editorial moment for an opener or a closer, taken when the photo is known to carry it.

DARK VARIANT • ANY LAYOUT CLASS

The dark modifier works on any layout.

Add *dark* alongside any class — palette remaps automatically

The token system handles dark without per-element overrides.

All colours reference CSS variables — `--bg`, `--text-heading`, `--text-body`, `--border` — that remap when `dark` is added. Cards, headings, body text, and borders all shift automatically. The spectrum bar is suppressed on dark slides.



Signal Pipeline · Reference Visualization

Every composition is palette-blind — add *dark* and the solid card, scrim, and matte all track the theme through the same tokens. The photo carries; the message stays legible.

The card stack renders cleanly on dark backgrounds.

Every card uses `--bg-alt` for fill and `--border` for the border — both remap in dark mode.

The accent left border uses `--accent` which is unchanged — the gold reads well against dark.

Body text shifts to `--text-body` which in dark mode is a warm light tone, not pure white.

Two-card layouts work equally well inverted to dark.

The framework introduces exactly one hard question — who owns the scoring weights, and what happens to every past decision the morning someone changes them. The next forty slides are a confident, well-resourced exercise in deferring the answer.

The pattern here is the same as any page of written argument — claim, then support. The dark palette does not change the information density or the reading rhythm.

The phase modifier renames the prefix word from STEP to PHASE.

PHASE 01

Architecture

The first phase scopes the technical surface — what we build, what we buy, what we defer. Output is an architecture decision record signed by the platform owner.

PHASE 02

Pilot

One internal team, one workload, one quarter. The phase ends when the integration is in production and the on-call rota covers it.

PHASE 03

Rollout

Five teams in two months. The phase ends when no team needs handholding and incident volume is at or below pre-rollout baseline.

MODIFIER · LIST-STEPS MILESTONE LETTERED

Modifiers compose: milestone renames the word, lettered swaps the format.

MILESTONE A

Scoring policy in production

The signed scoring policy runs end-to-end. The first calibrated brief lands in leadership's inbox, which we are calling general availability.



MILESTONE B

Per-team weights

One framework carries distinct scoring weights per team with no per-team fork. Recalibration is a single policy update. Each team will still want its own anyway.



MILESTONE C

Per-decision-class profiles

A scoring profile scoped to a single decision class takes minutes to author and a workshop series to agree on. Audit trails distinguish the classes nobody disputes.

Vertical stacks the steps as rows; the connector becomes a down-arrow.

STEP 01

Sense

Inputs are signals. Signals are observed, never invented. The first step is to write down what you see, not what you conclude.



STEP 02

Score

A signal becomes data once it carries a number. The score is calibrated against outcomes, not against intuition.



STEP 03

Decide

A decision is a signal plus a deadline. Without the deadline it is an opinion, not a decision. The retrospective closes the loop on the

Chosen flags the right-hand card as the winner.

QUARTERLY RE-LITIGATION

Every decision reopened from first principles each review. Close time is a function of seniority in the room, not evidence — average 4 hours, and the debate runs long precisely when the board joins.




CALIBRATED WEIGHTS

The decision resolves against logged weights and prior outcomes. Average 18 minutes, and the argument stops cascading into the next quarter — a property we discovered after building the loop, not before.

The right card carries an accent left-edge and accent-tinted background — the same visual contract used by other recommended-option cards.

Decision composes chosen + rejected with a labelled connector.



The left card is struck through to read as the option considered then dropped; the right card carries the chosen visual; the connector is amplified and labelled DECISION.

Vertical stacks the two cards; the arrow connector rotates 90°.

BEFORE — MANUAL RECALIBRATION

Operators book a recalibration window, freeze new decisions, swap weights, run a verification pass, lift the freeze. Average review pause 18 working hours; average post-mortem, ninety.

AFTER — VERSION-FLOOR RECALIBRATION

The calibration loop emits a new scoring policy with an incremented version. Teams pick it up on next refresh. No freeze, no cutover, no heroics worth a slide at all-hands.

Three switches the grid from 2 columns to 3 columns.

Signal

The observed input — a verbatim, a metric move, a competitor announcement. The unit of intake, and frequently confused with "things the VP heard at a conference," which score a 5 on relevance every time.

Decision

A signal plus a deadline, logged with rationale and predicted outcome. In theory traceable to its signals; in practice the signal is often "we discussed it at the offsite," logged after the fact, if at all.

Outcome

The observed result, compared at retrospective to what we predicted. The unit of calibration. Eighteen logged so far, against roughly three hundred forty that occurred.

Four switches to 4 columns; pair with compact for visual balance.

Sense

Signals are observed, never invented. Inputs are written down before they are interpreted.

Score

A signal becomes data once it carries a number. Calibration is against outcomes, not intuition.

Decide

A decision is a signal plus a deadline. Without a deadline it is an opinion.

Review

The retrospective closes the loop on the score that earned the decision. The model improves only here.

Horizontal flips cards-stack from a vertical stack to a row.

Claim


The framework buys calibrated prioritization with audit-grade decision custody — a sentence we will be repeating verbatim for two years.

Evidence

Six-month pilot, four product teams, decision close-time down to 18 min, calibration run once — measured by the team that predicted exactly this.

Implication

No vendor cutover. We keep funding the in-house build and ship org-wide enablement "next phase," as is tradition.



MODIFIER • IMAGE MIRROR

Mirror flips the image side — same vocabulary as split-panel watermark, compare-prose.

The image card moves from the right to the left, and the text padding swaps to match.

`mirror` is the cross-cutting side flag in the Lattice grammar; `![bg left]` is preserved as a backwards-compatible alias.

What this section covers

Mirror moves the dark accent panel to the right. The watermark, eyebrow, and section number all stay anchored to the panel's own box — only the column position flips.

1

Confidence

How many independent sources corroborate the signal. Ranges 1–5.

2

Recency

Time-decay applied from signal date to scoring date. Half-life is team-configurable.

3

Strategic Relevance

Manual score from the signal owner. Ranges 1–5. Requires justification above 4.

SECTION 02 • MIRROR

Section opener with the accent panel on the right.

Mirror composes with chosen — the accented card reads from the left.

THE CHOICE

With `mirror`, the chosen card now appears on the left visually. Use this when the surrounding deck reads right-to-left or when the chosen path needs to land first in the audience's scan path.



CONSIDERED ALTERNATIVE

Source order keeps this card first, so `chosen` rules continue to target the second card in the markdown. Mirror only flips the rendering; the editorial intent (left = considered, right = chosen) is preserved by reading order.

The below-note still appears under both cards after the hairline rule.

MODIFIER · DIVIDER NUMBERED

Numbered stamps an auto-counter in the top-right corner.

The CSS counter walks the whole deck once and increments on every *divider.numbered* slide. Authors do not number sections by hand — the layout does it.

MODIFIER · DIVIDER LIGHT NUMBERED

Each bookend layout owns its own counter.

The divider light counter is independent of the dark-divider counter, so a mid-deck light divider stamps **01** even when the dark dividers are already at **04**.

USE IT FOR MULTI-PART DECKS WHERE THE CLOSING SLIDE OF EACH PART SHOULD CARRY THE PART NUMBER.

CLOSING • NUMBERED

The closing series gets
its own auto-stamp too.

How we sort vendors against our two axes.

High coverage / Low cost

Sprig + Log — strongest coverage, lowest TCO, and a roadmap slide that is mostly our logo. Built by the evaluation team.

Productboard — narrower coverage but the cheapest tier, which is the only number procurement read.

High coverage / High cost

Notion build-out — full coverage in theory, premium maintenance, and seven slightly different versions of the framework.

Low coverage / Low cost

Chorus — cheap, but leaves three criteria uncovered, popular anyway.

Low coverage / High cost

none — and that is the signal.

We are building, not buying.

BUILD

Owns the scoring policy, the calibration loop, the timeline. And the pager.


WHY NOT BUY

Three vendors evaluated; none expose the calibration weights to the customer, and all three decks were the same deck.

WHY NOT DELAY

The competitive window closes in 18 months — two reorgs from now.

Decisions used to require a quarterly re-litigation.



The architecture change is the calibration loop — logged, weighted, time-bound scoring — not a meeting we ran until everyone stopped arguing.

How we make calls when the spec is silent.

01 We default to the choice that is cheaper to reverse.

02 We name the actor, never the system.

03 We write down the bet on the same slide as the choice.

What ships in each phase, by workstream.

WORKSTREAM	PHASE 01	PHASE 02	PHASE 03
Signal Intake	Connector v1	Multi-source dedupe	Anomaly auto-routing
Scoring	Equal-weights model	Per-team calibration	Per-decision profiles
Decision Log	Append-only schema	Outcome auto-pairing	Examiner export
Adoption	One pilot team	—	Org-wide enablement

The first column is sticky workstream label; phase columns carry numbered chrome; empty cells render as a thin dash.

Where we are against the targets we set ourselves.



94%

Signal-classification success

TARGET 99%, GAP IS "KNOWN ISSUE"



18 min

p99 decision close

TARGET 20 MIN, BEATING TARGET

18

Decisions logged

TARGET 340, GAP IS "CULTURAL"

1

Calibration cycles run

TARGET 6, GAP IS "STRUCTURAL"

What this deck covers, in order.

01 The Design — page 7

→ **The Phasing — page 18**

03 The Choices — page 26

04 Appendices — page 35

05 Closing — page 64

Who owns each part of the framework lifecycle.

Signal custody

Signal owner

Runs intake quality. Never tunes the weights — only picks which signals surface.

Policy

Framework operator

Owns scoring policy, calibration cadence, and the rollback playbook nobody has run. One person.

Consumption

Product team

Holds scoring profiles; runs intake and decision-logging; finds the bugs first.

Oversight

Auditor

Reads the audit trail, cannot edit weights, and is the only role anyone fears — having read it once.

What this section will tell you, in five lines.

- 01

The framework buys calibrated prioritization with audit-grade decision custody. → slide 8
- 02

Recalibration is a version-floor increment, not a coordinated freeze, not a war room. → slide 12
- 03

Per-team weights make recalibration a single policy update. → slide 18
- 04

Phase 1 ships the architecture, Phase 2 ships the operations, Phase 3 ships the apology. → slide 22
- 05

Five questions stay open until Phase 1 is forced to close them on the record. → slide 27

Compact tightens the spacing scale ~25 %, end-to-end.

What changes

`--sp-xs` through `--sp-2xl` shrink. Card gaps, list gutters, and section padding follow because every layout reads them via `var()`.

What does not change

Type ramp, palette, and chrome reservation (header / footer / pagination) are untouched. Compact is a density flag, not a different layout.

When to reach for it

You have one more card than fits, or your prose runs the section by 1-2 lines, or you want a denser visual rhythm without rewriting copy.

Composition

`compact` composes with `dark`, `accent`, and any layout where density makes sense. It is silently incompatible with `title`, `divider`, and `image-full`.

Loose is the inverse — more breathing room, same layout machinery.

The spacing scale grows ~25 % rather than shrinks. Sections that already look generous become luxurious; sections that look cramped become balanced. Reach for **loose** when the slide carries a single editorial point and you want the page to feel deliberately quiet — values pages, declarative principles, the closing line of an argument.

The discipline is the same as **compact** from the other side: do not change the type ramp, do not change the chrome, do not change the layout. Only

Headings gain a closing period automatically.

Authors who prefer sentence-style heading punctuation can set `class:` `with-period` in front matter once and stop thinking about it. The transform appends a period to any heading that does not already end with terminal punctuation — `.` `!` `?` `:` `...` — so mixed slides are safe.

The mirror modifier is `no-period`, which strips trailing periods instead.

Both are deck-wide opt-ins via the global `class:` front-matter key; per-slide override with `←!— _class: with-period →` works too.

Authors typed this heading with a period. It is gone

Some teams author headings with periods out of habit, then strip them in review. `class: no-period` automates the strip so the source can stay as written and the output stays clean.

Only a literal trailing `.` is removed — `!`, `?`, `:`, and `...` pass through untouched. Combine with any layout class; the modifier composes cleanly because it operates on the heading text alone and touches no structural chrome.

TREATMENT LIBRARY • ANY LAYOUT CLASS

`tint-*` and `mark-*` classes add peripheral accents from the active palette

Add a tint or mark class alongside any layout class — gradient wash or SVG mark, light canvas or dark, single pattern or layered pair.

A radial glow anchored at the corner fades before reaching the content zone

`tint-corner at-tl` places an elliptical accent at the top-left — 12% opacity at the corner, transparent before mid-slide. The four `at-*` placements share the same weight and fade profile; only the anchor differs.

All gradients use `color-mix(in srgb, var(--accent) 12%, transparent)`. Switching palette or adding the `dark` modifier remaps the accent automatically — no per-pattern overrides.

SVG accent marks are painted through a mask in the active accent colour

`mark-orbit` places concentric rings and satellite dots in the bottom-right corner. The shapes render via `::before` + `mask-image`: the SVG defines the alpha channel (white = opaque, transparent = hidden) and the paint colour is `color-mix(in srgb, var(--accent) 28%, transparent)` — resolved from the theme at render time. Same class, light canvas or dark, the shapes are always visible and always on-brand.



One class from each slot layers without conflict

Every tint class (and `mark-seeds`) writes to either `--_bg-radial` or `--_bg-linear`. A compositor rule assembles both slots into a single `background-image` with two live layers. Stack one class from each column and both render:

- `tint-vignette` — radial slot — accent-tinted perimeter, open center
- `tint-edge at-right` — linear slot — wash bleeding in from the right edge

The SVG mark patterns follow the same rule: their atmospheric haze writes to its slot, and the `::before` shapes compose on top independently.

CHART LAYOUTS · GANTT + KANBAN

Timeline bars and board columns from a two-level list

2026 Q1 .. 2026 Q4

Feature delivery by workstream



Where Phase 2 work stands today

BACKLOG

API contract review

Load-test harness

IN PROGRESS

SDK v2 alpha in-progress

Onboarding redesign in-progress

REVIEW

Token migration spec review

DONE

Scope sign-off done

Design freeze done

SPLIT LAYOUTS · FIVE VARIANTS

A structured left panel drives five two-column layouts

Q2 PERFORMANCE REVIEW

Enterprise revenue stalled in Q2

Three structural factors explain 90% of the shortfall — all addressable before Q4 close.

Split — brief • split-panel

Renewal pricing complexity is driving churn at the segment ceiling

Four accounts totaling \$2.1M ARR declined renewal. Win/loss interviews point to a quote-to-contract gap, not value perception.

Pipeline conversion dropped 11 pp below Q1 — legal review is the chokepoint

Contract length increased 18 days on average. Root cause is a security addendum introduced in March.

Competitive displacement accelerated in the \$80–200K ACV tier

Seven losses to a single competitor. Time-to-value gap is the exposure.

NET REVENUE RETENTION

114%

Measured across all customers active for 12+ months, March 31 cohort.

Existing customers are growing faster than we lose them

At 114%, every churned dollar is offset by \$1.14 in expansion. The base compounds without new-logo dependency.

Expansion is concentrated — three segments drive 80% of the gain

Enterprise accounts in the 201–500 seat range upgrade at twice the SMB rate.

Sustained above 110%, this unlocks a capital-efficient growth path

NRR above 110% meets the investor threshold for venture-category efficiency.

02

Discovery & Scoping

Four weeks. Shared definition of the problem before any solution work begins.

- 1

Stakeholder Interviews
Eight cross-functional conversations. Open questions only — listening for friction, not confirming assumptions.
- 2

Current-State Audit
System inventory, workflow documentation, and data quality review.
- 3

Problem Framing Workshop
Half-day session to align on root cause. Output is a ranked problem statement the team signs off on.
- 4

Scope Confirmation
Written sign-off on what is in, what is out, what requires a separate decision.

DECISION REQUIRED

Build the data layer or buy it?

Both paths are viable. The difference is where we spend the next 18 months.

BUILD IN-HOUSE

- Full control over schema and roadmap
- 2–3 engineer-quarters to reach feature parity
- Ongoing maintenance burden stays internal

BUY + CONFIGURE



- Ship in 6 weeks, not 9 months
- Engineering capacity redirects to product-layer features
- Exit risk manageable — data export contractually guaranteed

RECOMMENDATION

Buy the infrastructure. Build the differentiation. Revisit in 24 months.



The best product does not win. The most understood product does.

Morgan Chase · Head of Product, Vercel, 2024

Split — statement · split-panel pullquote

Clarity is a product decision, not a marketing one

If a prospect cannot articulate our value in one sentence, the product has a communication architecture problem.

Onboarding is the product's first argument for itself

The moment a user first succeeds defines their frame for everything that follows.

Understanding, not delight, is the retention driver at scale

Users who understand the system's logic stay through friction. Build for comprehension.

The default top border is a spectrum gradient — a system signal that the page is part of a wider deck. The `ACCENT` modifier swaps that stripe for one solid colour and tints the slide heading. Use it when one slide carries the editorial weight of a section and you want the visual chrome to say so.

MODIFIER • ACCENT

Accent replaces the rainbow stripe with a single editorial colour.

It composes with `dark`: on the dark canvas the spectrum top-stripe is suppressed entirely, so `accent.dark` restores a solid accent stripe in its place — preserving the visual signal across both canvases.